

# Searching With No Flashlight

An overview of derivative-free optimization

# What is a Derivative-Free Algorithm?

## **Derivative-free (non-gradient) algorithm:**

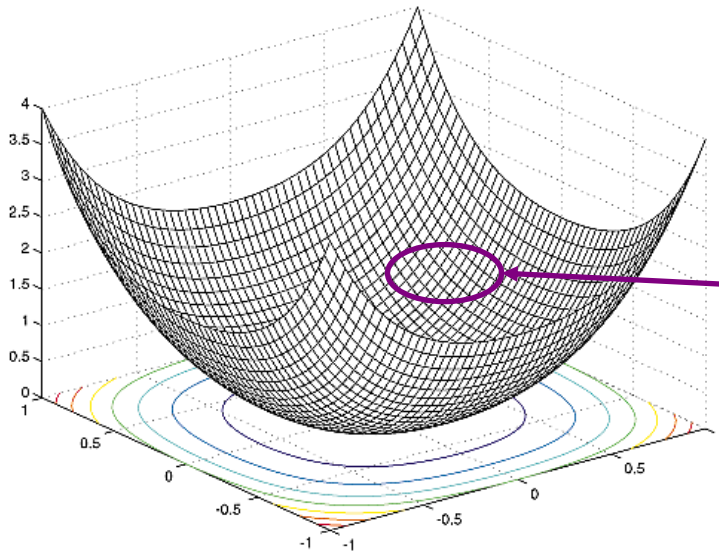
- No gradient information necessary
- “Smart” method of searching design space based upon some heuristics

## **Outline:**

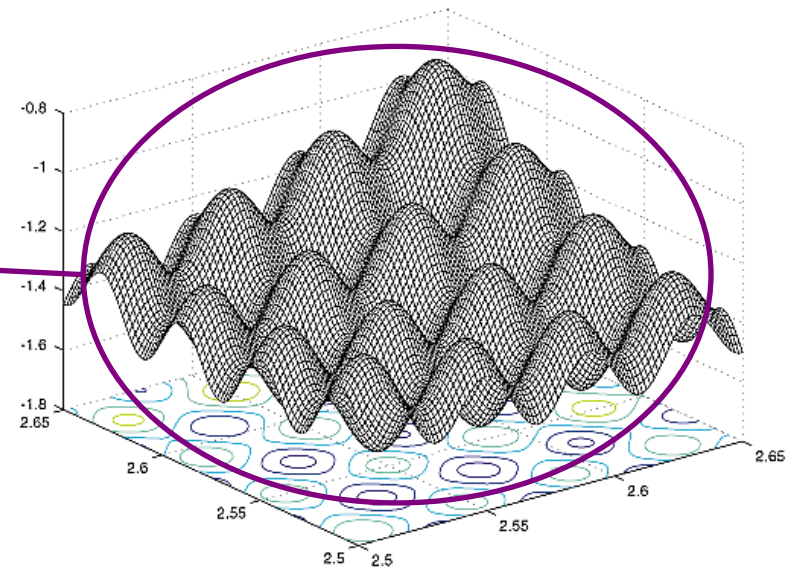
- Why use derivative-free algorithms? And why not?
- Review of existing algorithms

# Why Derivative-Free Algorithms? (1)

- Expensive function evaluation
- Noisy function evaluation



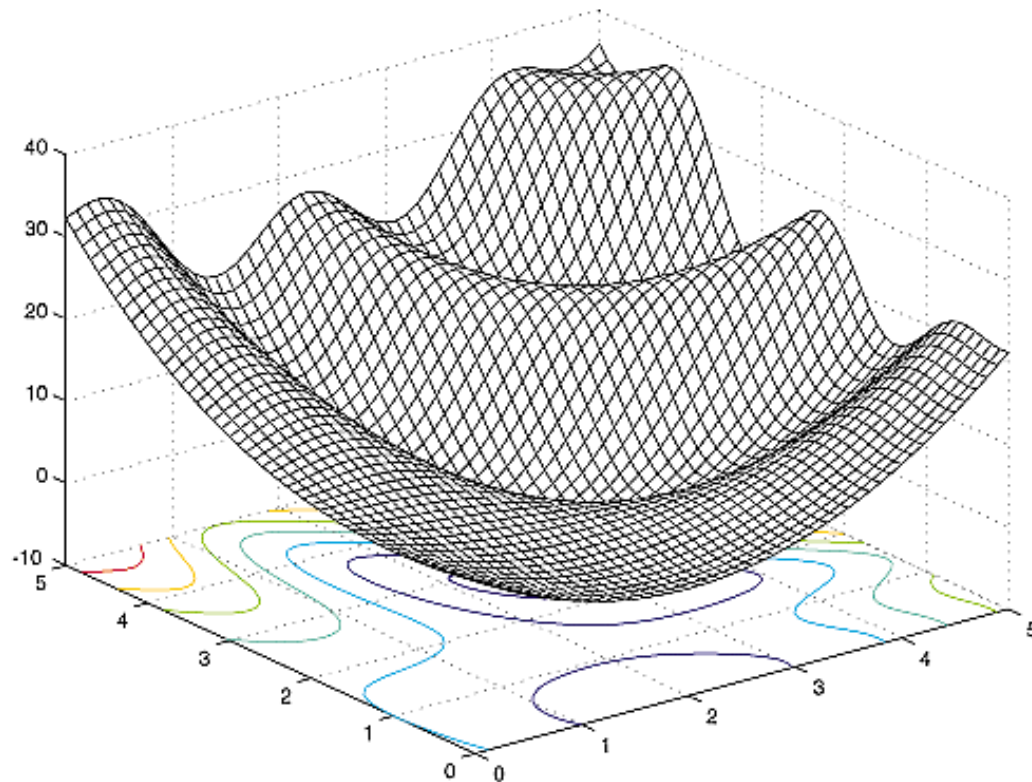
unimodal function



numerical noise

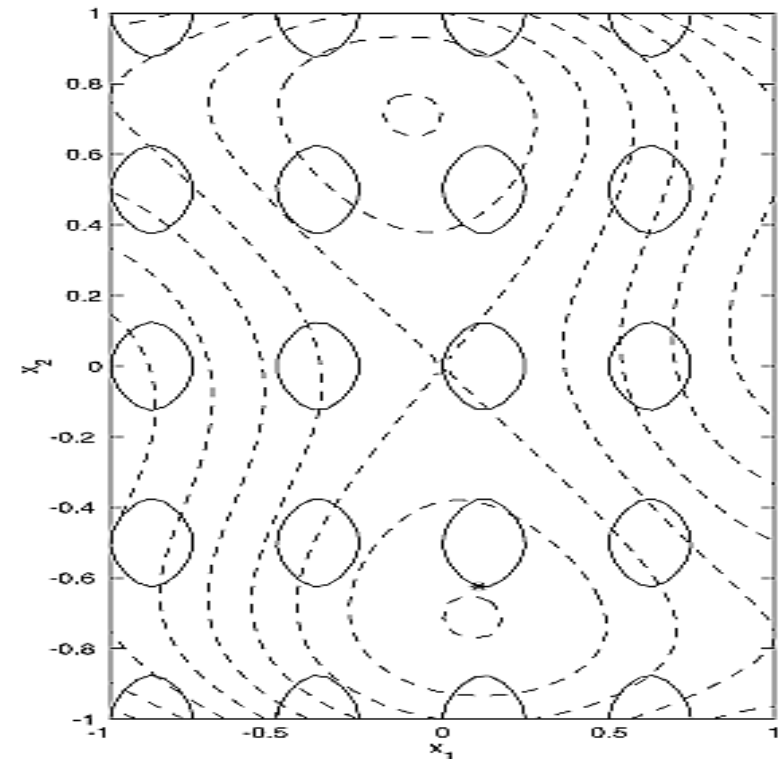
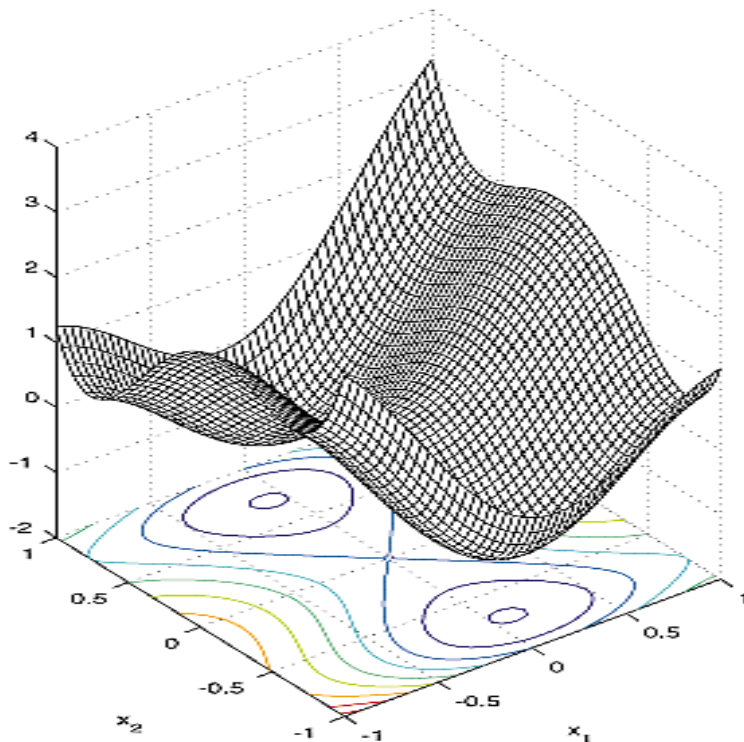
# Why Derivative-Free Algorithms? (2)

- Multiple optima exist



# Why Derivative-Free Algorithms? (3)

- Disconnected feasible regions
- Difficulty in finding feasible points



disconnected feasible region

# Why Derivative-Free Algorithms? (4)

- **Discrete choice variables / combinatorial problems**
  - Material selection
  - Component selection
  - Routing problems
- **Integer Variables**

# Why NOT Derivative-Free Algorithms?

## Disadvantages

- Slow to converge
- Usually no guarantee of optimality
- Often require tuning of many algorithm parameters
- Constraint handling often through penalty functions
  - No guarantee of feasibility
  - Equality constraints are more difficult

# Classes of Derivative-Free Algorithms

## **Stochastic**

Search depends on probability/random number generation;  
Each run of algorithm will take different search path and may find different “best point”

## **Deterministic**

Search follows distinct path (dependent on starting point, if specified); Each run of algorithm will have same result



# Existing Derivative-Free Algorithms

## Stochastic methods

- Simulated annealing
- Genetic algorithms
- Particle swarm

## Deterministic methods

- DIRECT
- Multilevel coordinate search (MCS)
- Efficient global optimization (EGO)
- NOMAD (hybrid method)

**and MANY others...**

# Survey of Derivative-Free Algorithms

**Exhaustive survey** by Rios and Sahinidic:

- 22 algorithms considered;
- On over 500 problems (convex/nonconvex + smooth/nonsmooth) with bounds only;
- With #variable from 1 to 30;
- Limit of 2500 iterations and 600 CPU seconds.

## Conclusions

- There always exist a few problems that a certain solver has the best solution quality.

[http://egon.cheme.cmu.edu/ewocp/docs/SahinidisEWO\\_DFO2010.pdf](http://egon.cheme.cmu.edu/ewocp/docs/SahinidisEWO_DFO2010.pdf)

# Topics

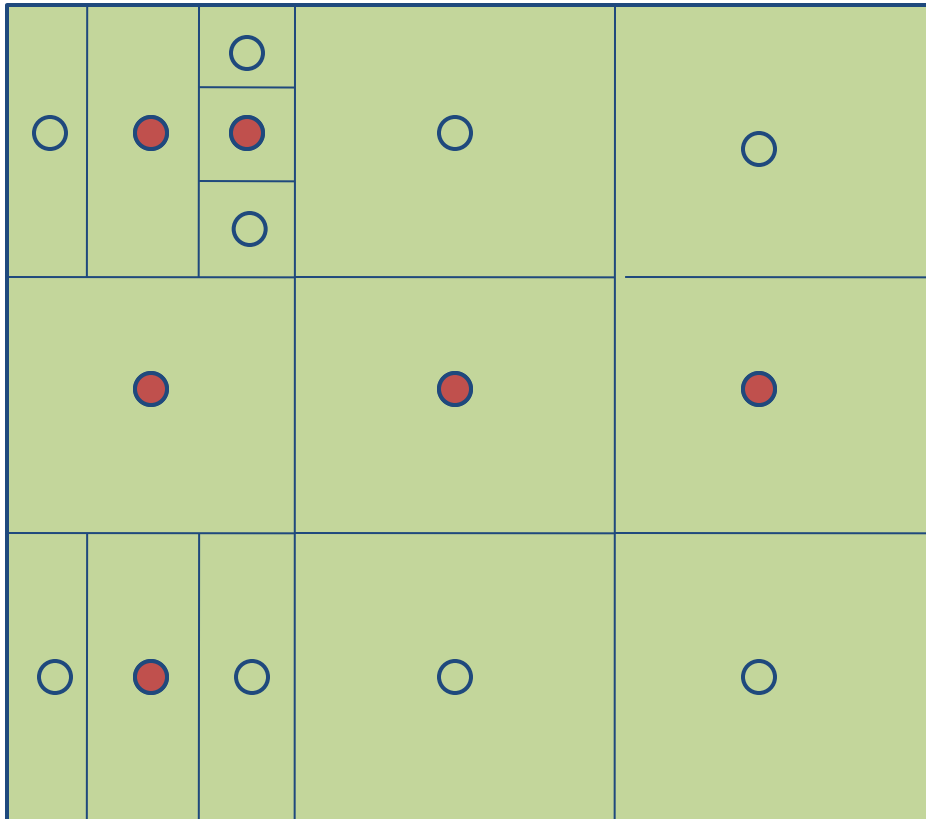
- **DIRECT**
- **Simulated annealing**
  - **Quantum annealing**
- **Genetic algorithm**
  - **CMA-ES**
- **Efficient global optimization (EGO)**
- **Pattern Search**
  - **NORMAD**

# DIRECT Overview

DIRECT stands for “Divided Rectangles”

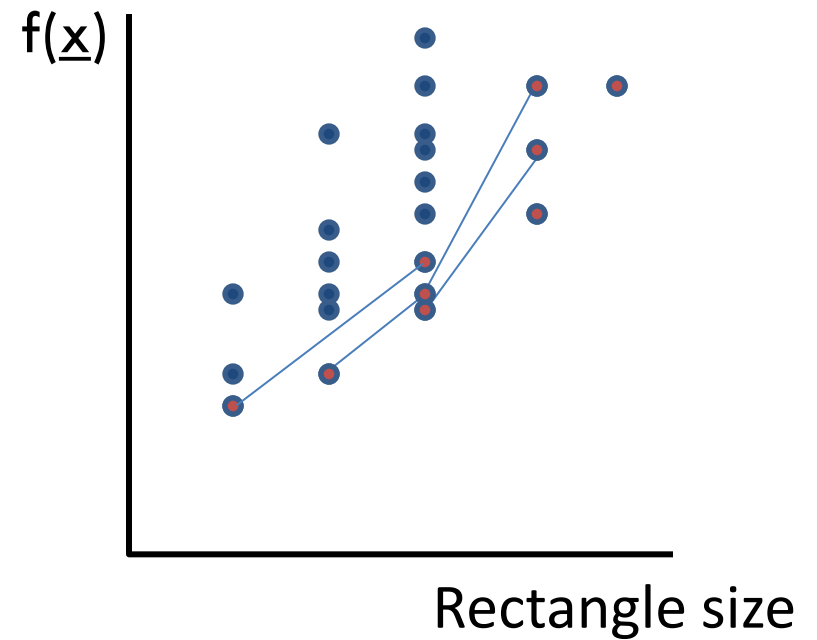
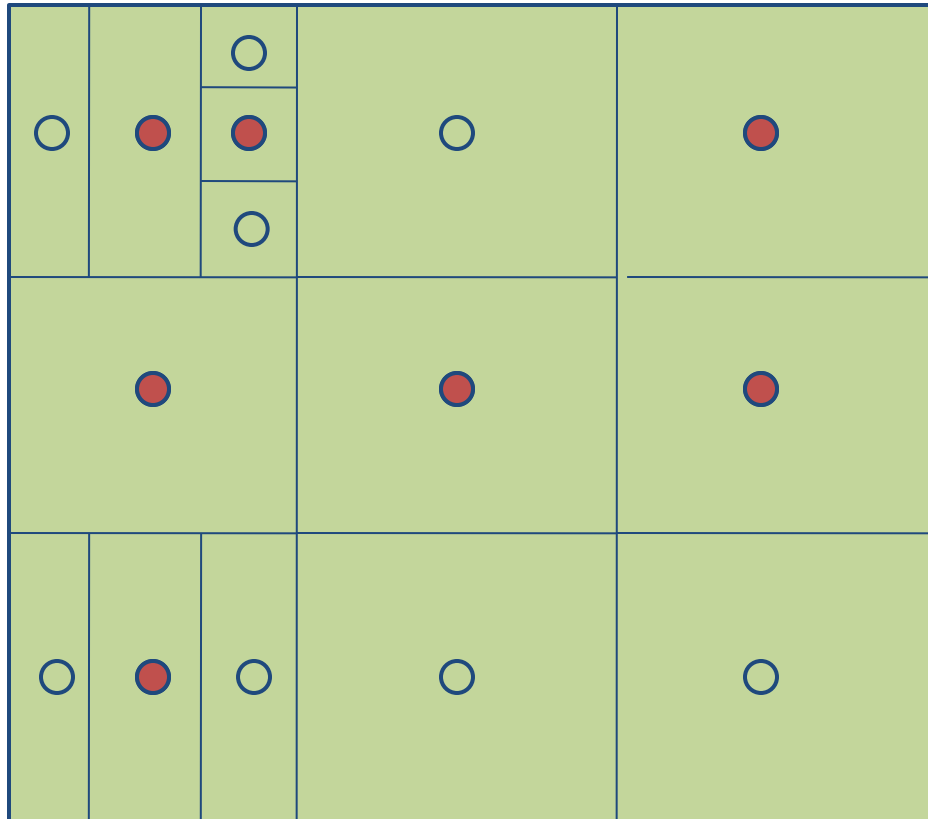
- Whole design space is sub-divided into rectangles;
- The “best” and “largest” rectangles are further divided.

# DIRECT with 2 Variables

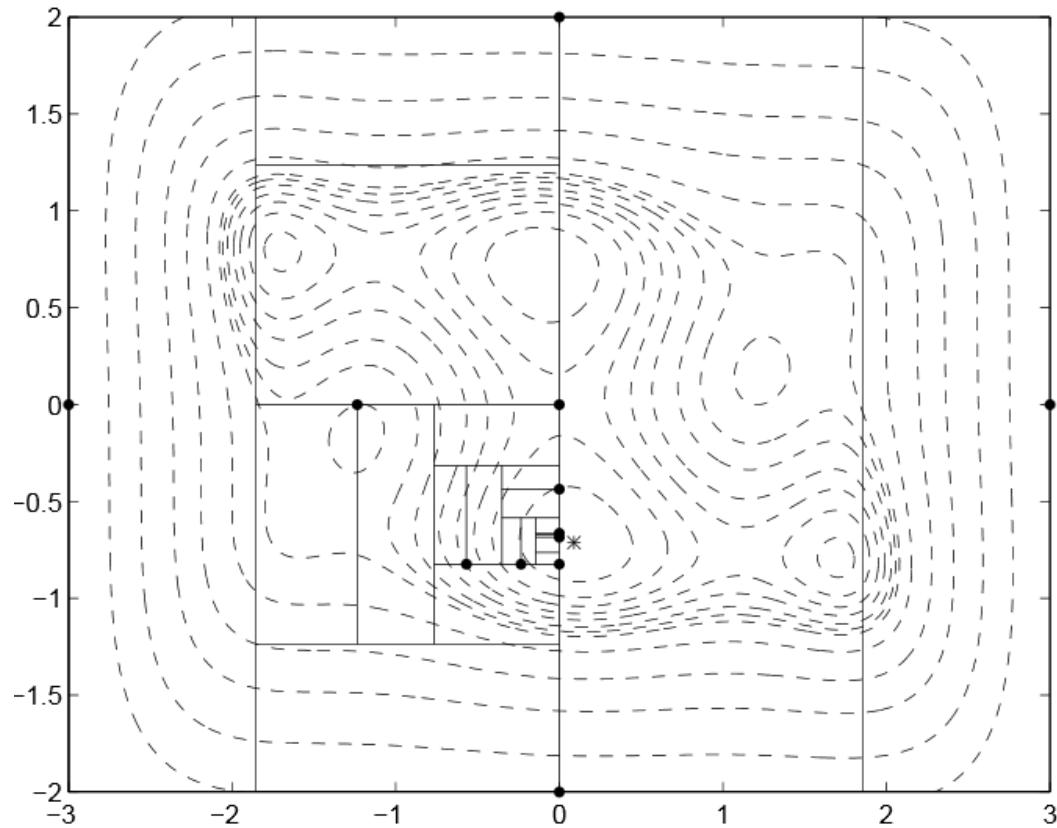


1. Sample center of design space
2. Select best candidate rectangles and divide into thirds along their longest dimensions
3. Best candidate rectangles based upon:
  - best  $f(x)$
  - lowest constraint violation
  - size of rectangle
4. Iterate until max. number of function calls

# DIRECT with 2 Variables



# Multilevel Coordinate Search (MCS)



Extension of DIRECT to have “basepoints” not in the center of boxes

# DIRECT Pros/Cons

## ■ Advantages

- Global and local search balance
- Deterministic, has the ability to be restarted where it left off
- No parameters to tune
- Can handle integer variables

## ■ Disadvantages

- Dimensionality: For problems of 10 variables or larger, DIRECT has difficulties because of having to divide along each dimension
- Slow local convergence
- Cannot handle equality constraints



# Simulated Annealing Overview



**Idea:** Simulate the cooling a metal to find the “strongest” configuration of atoms

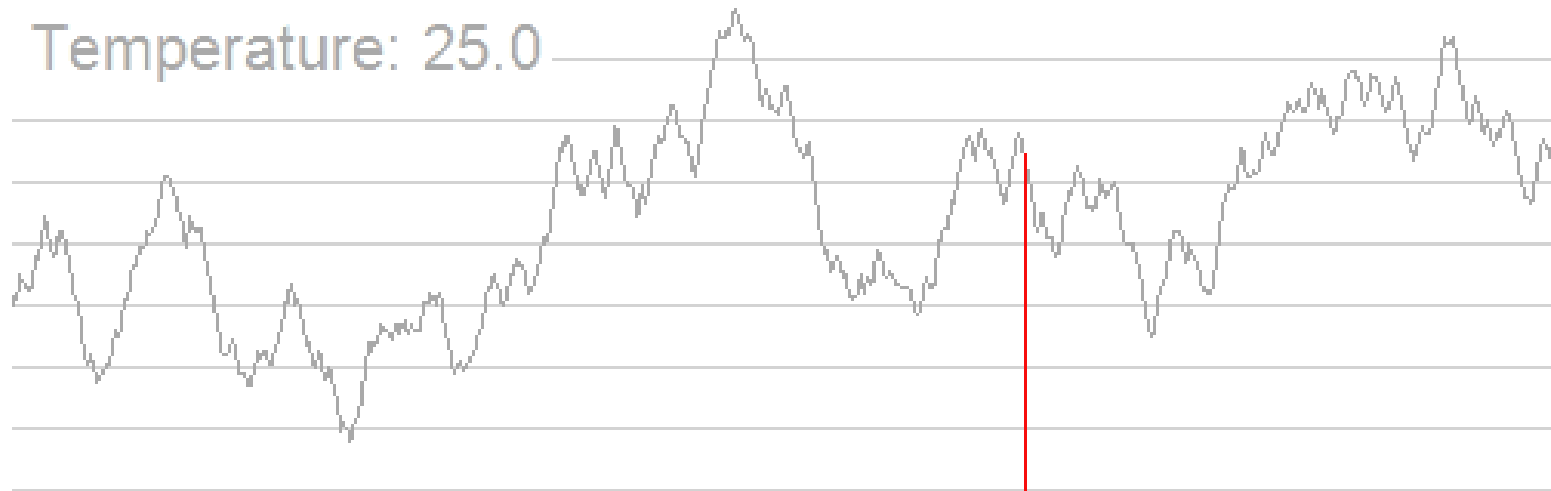
# Simulated Annealing - Algorithm

- Cooling of metals: want to find lowest energy state
- Performs random search with some probability of accepting a worse point (to get out of local minima)

$$\text{Prob}(\mathbf{x} \leftarrow \mathbf{y}) = \begin{cases} 1 & \text{if } \Delta f < 0 \text{ (better: downhill)} \\ \exp(-\frac{\Delta f}{t}) & \text{if } \Delta f \geq 0 \text{ (worse: uphill)} \end{cases}$$

- $t$  is the temperature at the current iteration.  $t$  decreases along the iteration number.

# Simulated Annealing - Demo



# Simulated Annealing - Constraints

Penalty function:

$$\min f_P(\bar{x}, \text{Penalty}) = f(\bar{x}) + \sum_{i=1}^m w_i \cdot (\max(0, g_i(\bar{x})))^2$$

- Most common is quadratic penalty function, though others are possible
- No guarantee of feasibility
- For equality constraints, can use two inequalities for upper and lower bounds
- Scaling of constraints and objective is ESSENTIAL to ensure feasibility with reasonable descent

# Simulated Annealing - Convergence

Proved by Geman brothers:

If the temperature drops by the following form:

$$T(t) = \frac{cN}{\log t},$$

where  $c$  is a problem dependent constant,  $N$  is the problem size (number of variables), then SA is guaranteed to find the global solution in infinite time limit.

# Simulated Annealing – Pros/Cons

## Advantages:

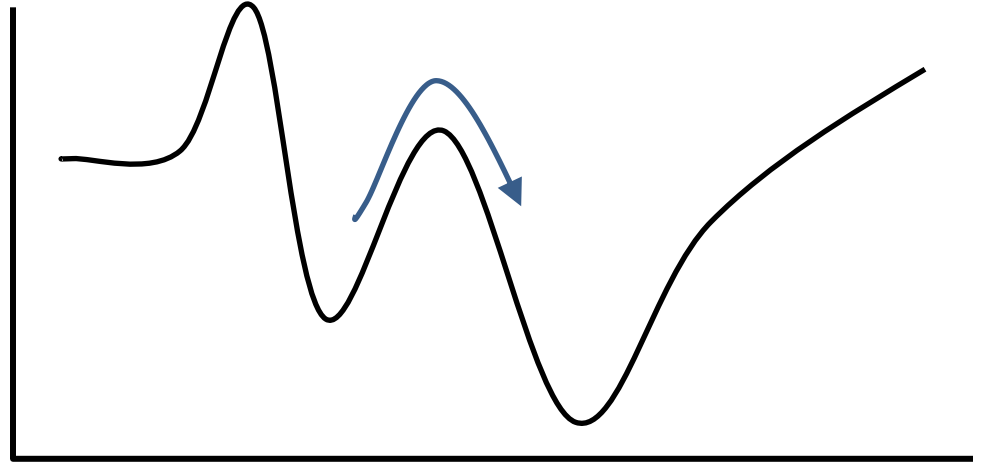
- Doesn't need to systematically cover space—better efficiency for large-dimension problems

## Disadvantages:

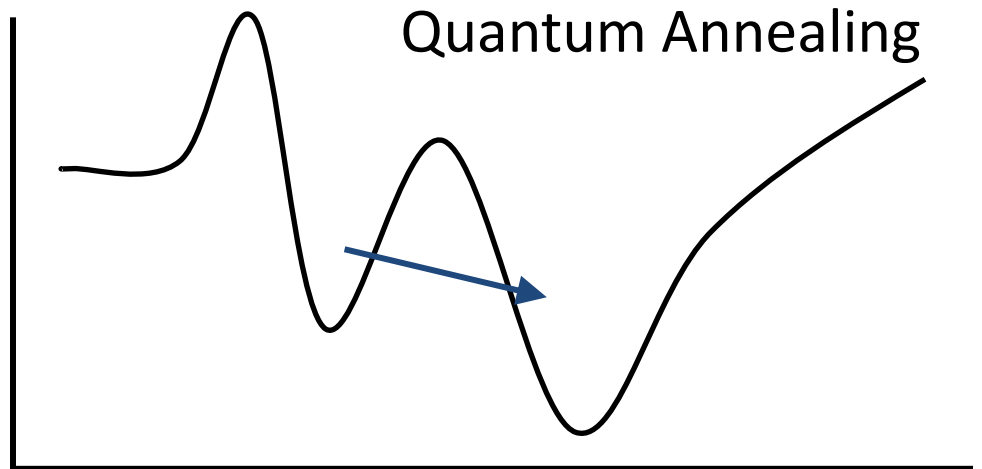
- Doesn't always cover the design space (quasi-global)
- Dependent on starting point
- Random directional search not very “smart”
  - Can repeat areas already searched
  - Can require large # of function calls
- Many parameters to tune – algorithm performance is dependent on these parameters
  - Penalty weights
  - Temperature cooling schedule

# From Simulated Annealing to Quantum Annealing

Simulated Annealing – Search  
by **thermal** fluctuation

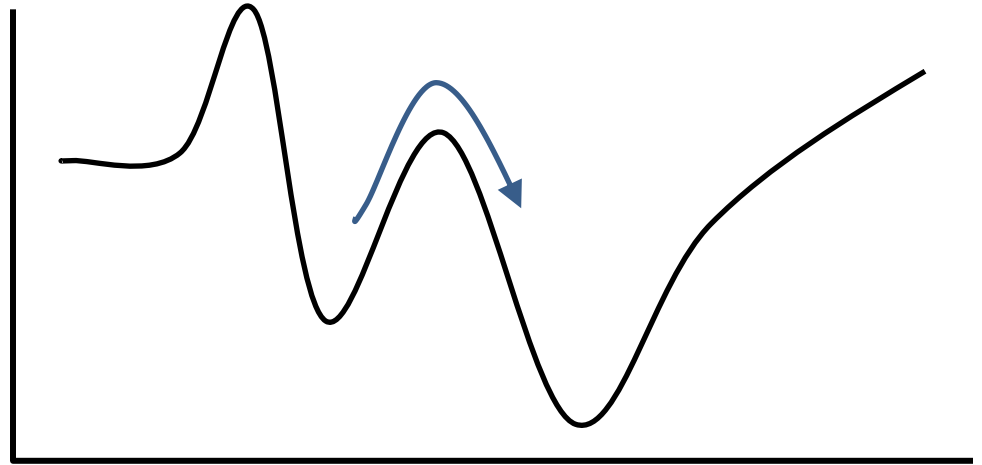


Quantum Annealing  
(Adiabatic Quantum  
Optimization) – Search by  
**quantum** fluctuation

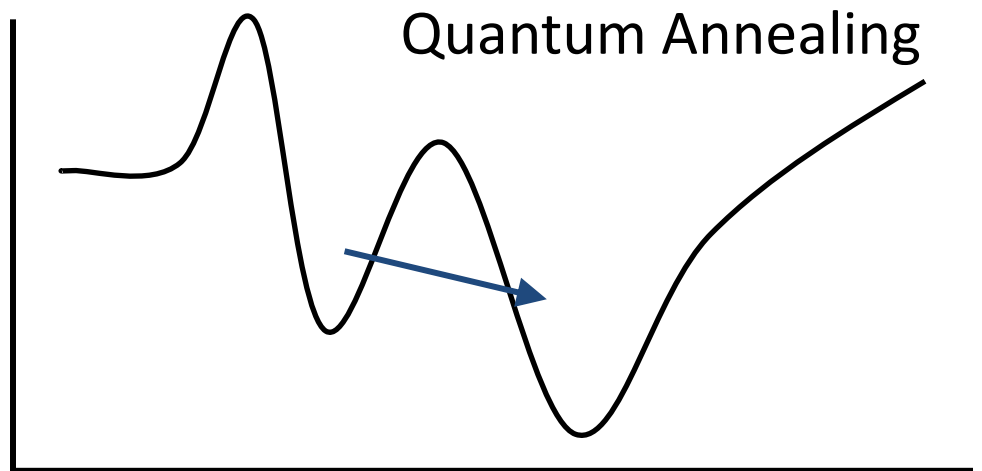


# From Simulated Annealing to Quantum Annealing

Simulated Annealing – probability to escape depends on the temperature and the barrier height (energy difference)

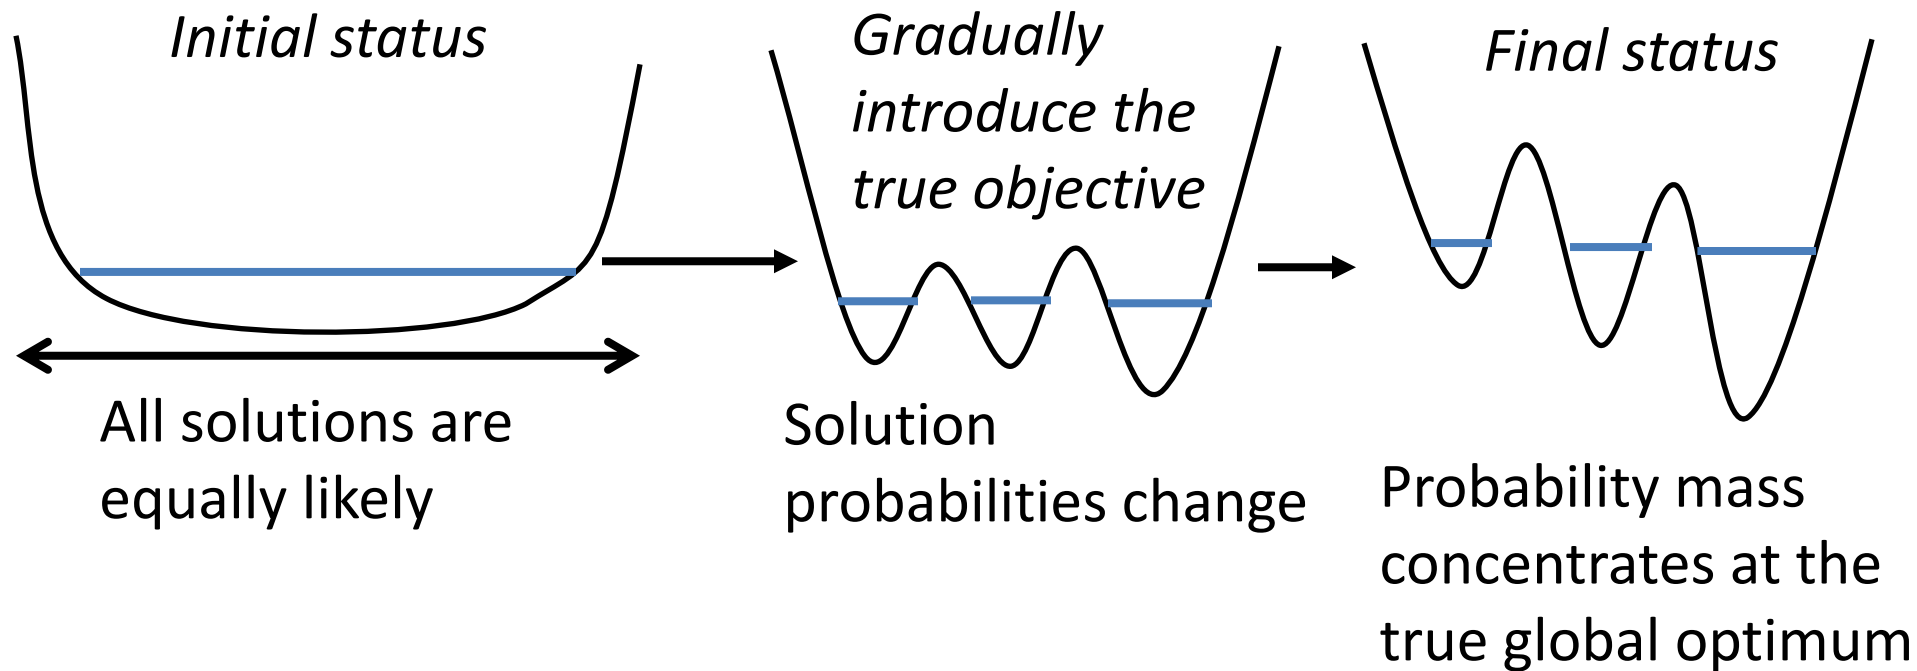


Quantum Annealing – probability to escape depends on the **quantum tunneling width** and the barrier width





# Quantum Annealing - Idea

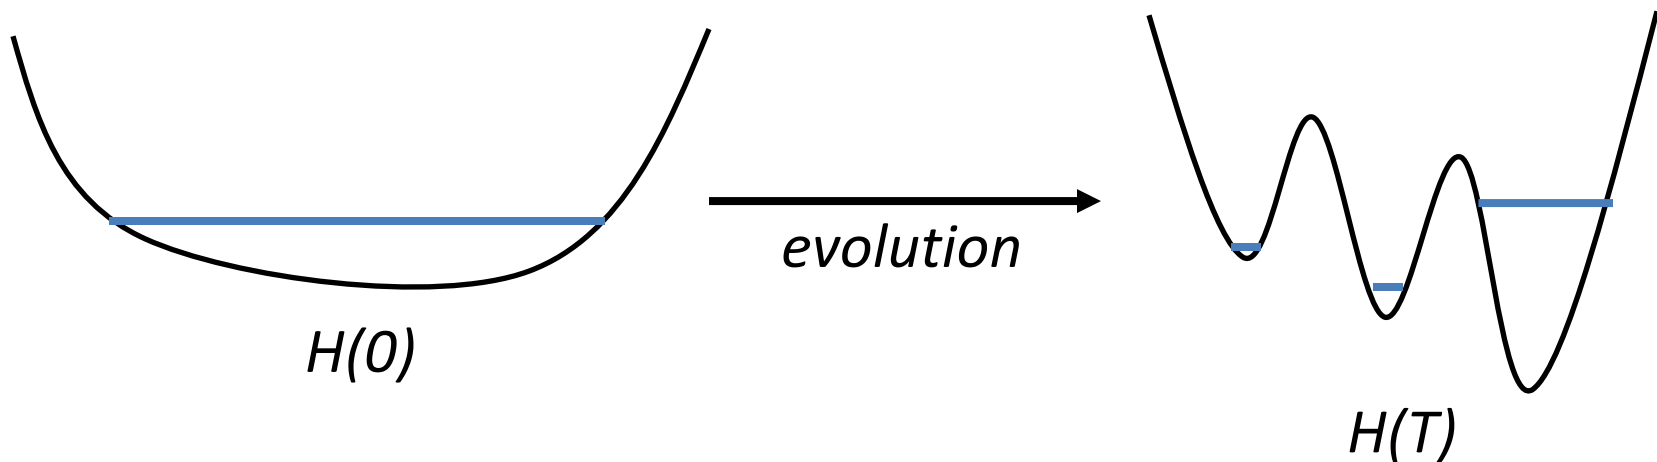


# The Adiabatic Theorem

Slowly varying Hamiltonian (objective) for evolution from  $t = 0$  to  $t = T$

Provided  $T$  is “large enough”, a quantum system starting in the ground state of  $H(0)$  evolves to the ground state of  $H(T)$

Large enough  $T$ :  $T = O(g^{-2})$ , where  $g$  is the minimum gap between ground state and first excited state of  $H(t)$

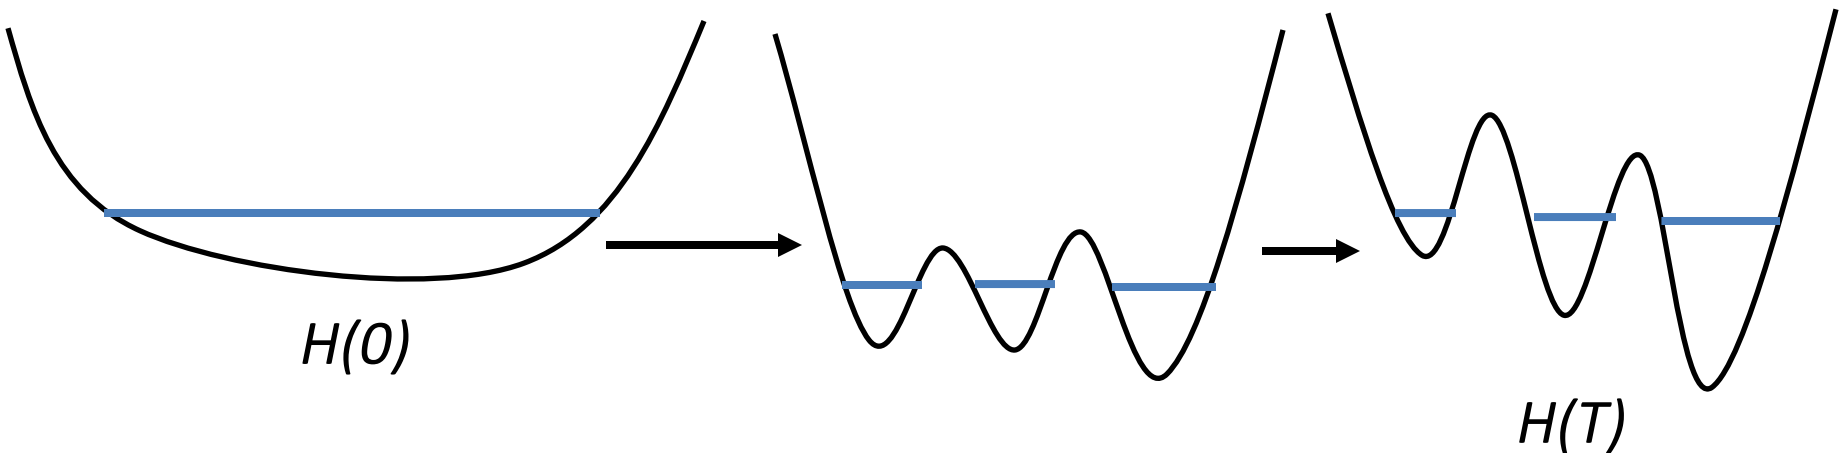


# The Adiabatic Theorem

Slowly varying Hamiltonian (objective) for evolution from  $t = 0$  to  $t = T$

Provided  $T$  is “large enough”, a quantum system starting in the ground state of  $H(0)$  evolves to the ground state of  $H(T)$

Large enough  $T$ :  $T = O(g^{-2})$ , where  $g$  is the minimum gap between ground state and first excited state of  $H(t)$



# Quantum Annealing – Convergence

SA convergence:

$$t = \exp\left(\frac{cN}{\delta}\right)$$

QA convergence (faster than SA for small  $\delta$ ):

$$t = \exp\left(\frac{N \ln \delta}{2c'}\right)$$

# Quantum Annealing – Implementation

D-wave QA solves the Ising function (Quadratic Unconstrained Binary Optimization)

$$E(x_1, \dots, x_N) = \sum_{i=1}^N h_i x_i + \sum_{i < j=1}^N J_{ij} x_i x_j$$

Can be applied to many machine learning problems (why?)

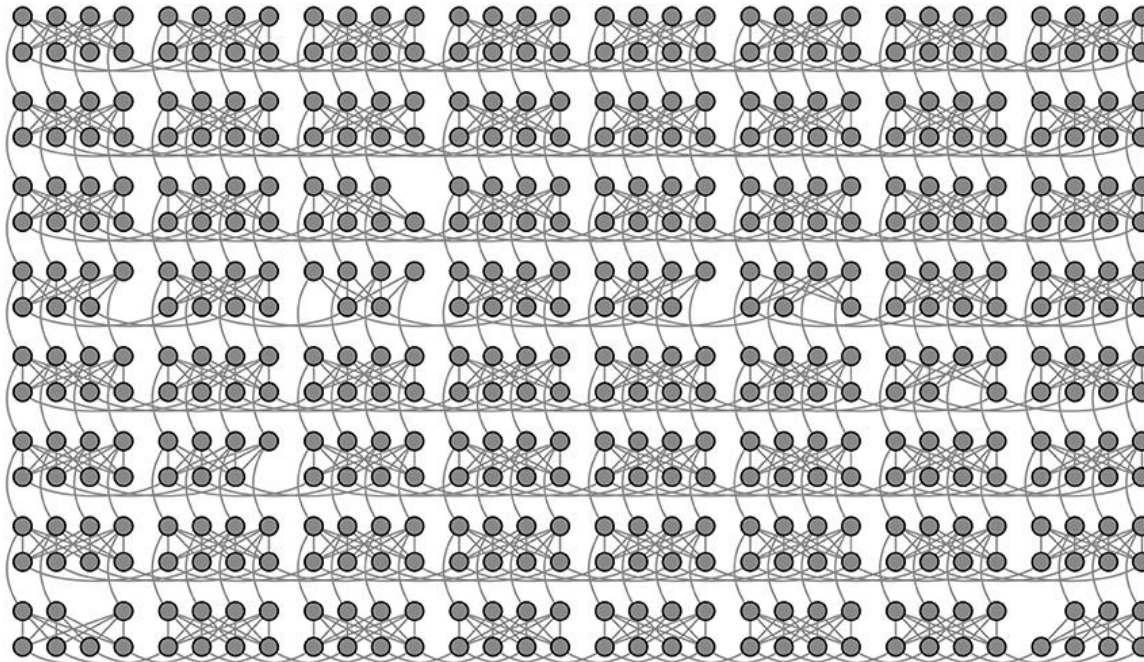
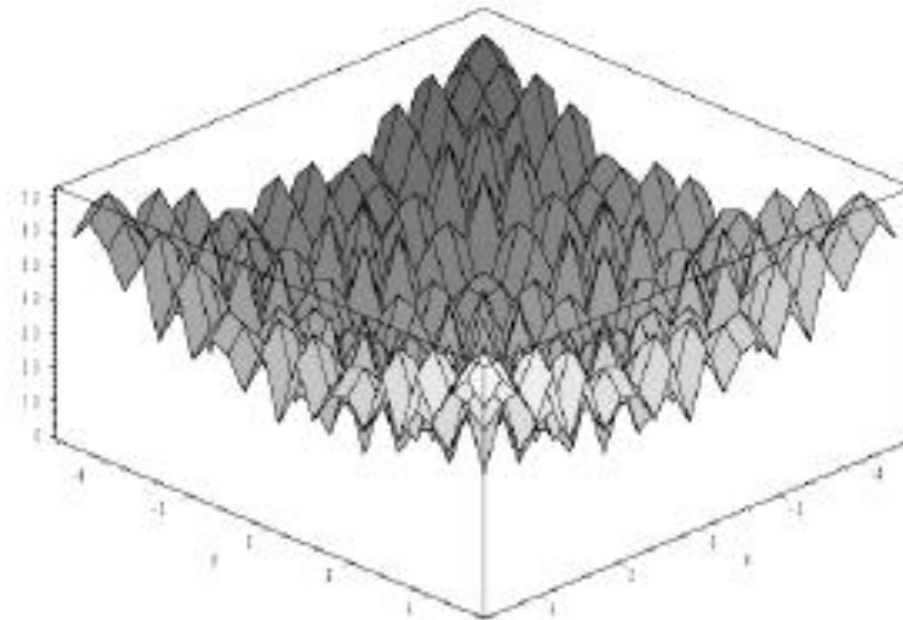


Figure from: [http://www.frontiersin.org/files/Articles/107579/fphy-02-00056-HTML/image\\_m/fphy-02-00056-g001.jpg](http://www.frontiersin.org/files/Articles/107579/fphy-02-00056-HTML/image_m/fphy-02-00056-g001.jpg)

# Quantum Annealing – Summary

QA has an advantage on problems “with **thin barriers** that separate very **deep chasms** between local minima”.



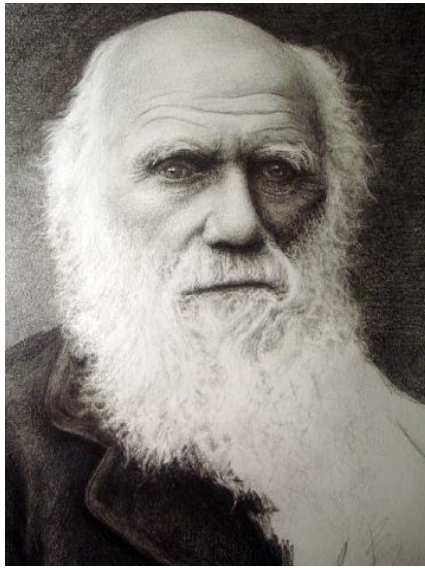
2D Rastrigin's function

More info:

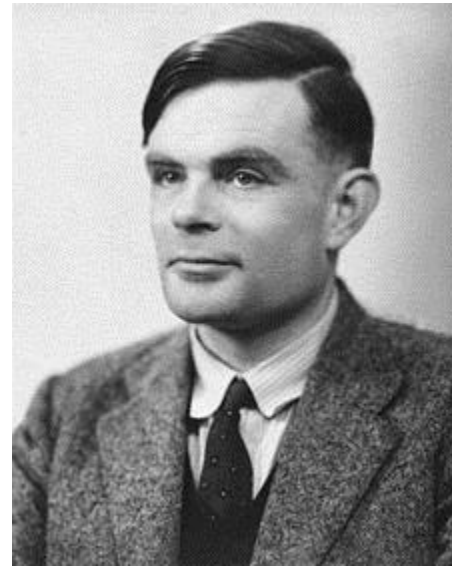
[http://www.stat.phys.titech.ac.jp/~nishimori/QA/q-annealing\\_e.html](http://www.stat.phys.titech.ac.jp/~nishimori/QA/q-annealing_e.html)

<https://plus.google.com/+QuantumAILab/posts>

# Genetic Algorithm



Charles Darwin

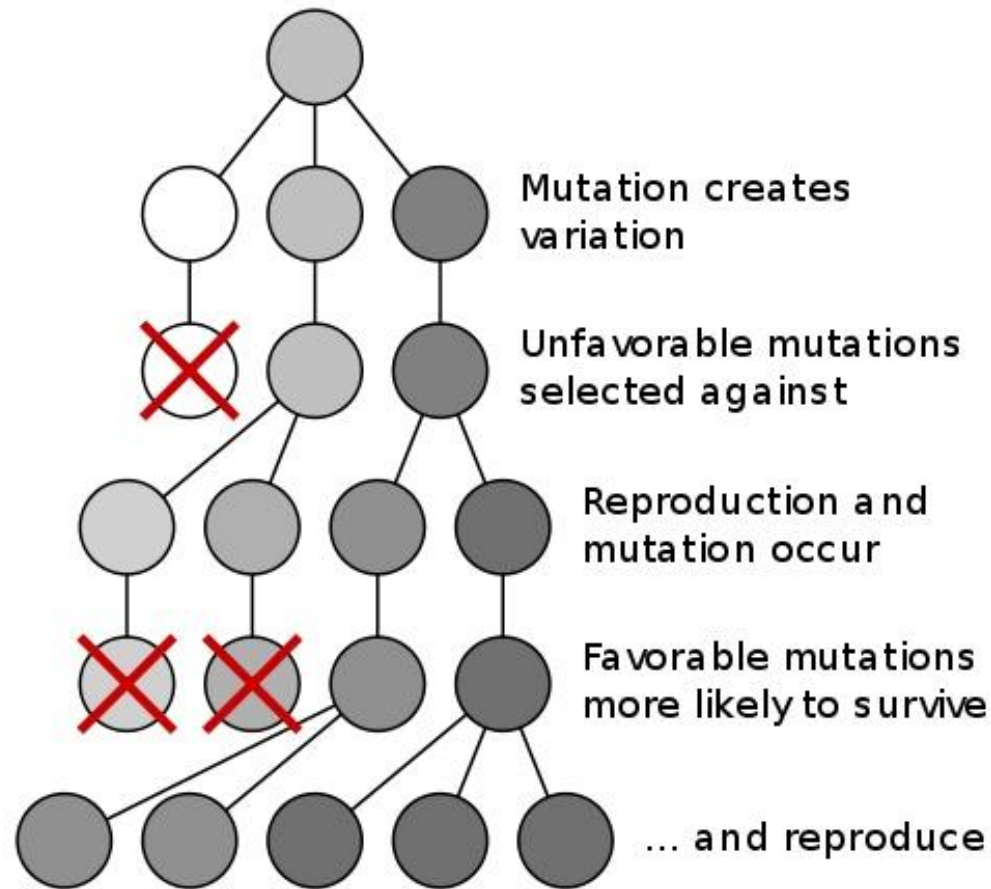


Alan Turing

Introduced by Turing, popularized in the 1980's

One of most popular nongradient methods

# Genetic Algorithm - Idea



[http://media.tumblr.com/036028e2fcf2db98eb94cb3dff0aca5f/tumblr\\_inline\\_mmzvcqpVQ71qz4rgp.jpg](http://media.tumblr.com/036028e2fcf2db98eb94cb3dff0aca5f/tumblr_inline_mmzvcqpVQ71qz4rgp.jpg)

**Idea:** Only allow the “fittest” designs move their DNA to the next generation

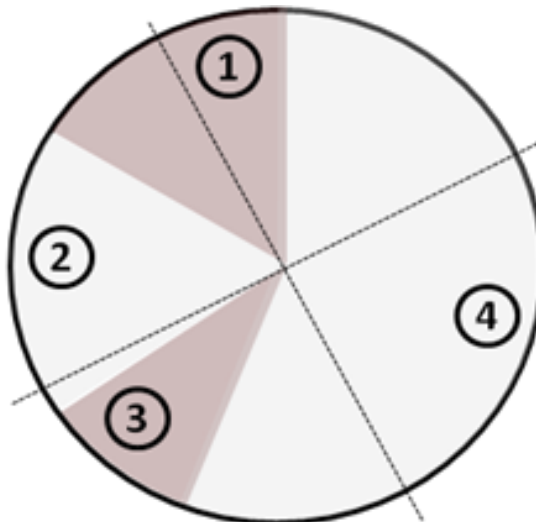


# Genetic Algorithm Overview

Starting with a population of random points in the feasible set, produce a new population of better points by *parent selection*, *crossover*, and *mutation*, until some conditions are satisfied.

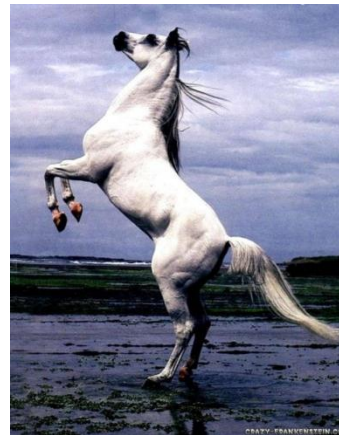
# GA - Parent Selection

- Many methods: roulette wheel, tournament, elitism, etc.
- Roulette wheel selection
  - Better individuals get larger portion of wheel
  - Random selection from wheel determines parents of next generation



# GA - Parent Selection

- Many methods: roulette wheel, tournament, elitism, etc.
- Tournament selection
  - Randomly pick  $k$  chromosomes from the population
  - Pick the best one out of the subset
  - Iterate until all parents are picked



**Each time pick three and compete**

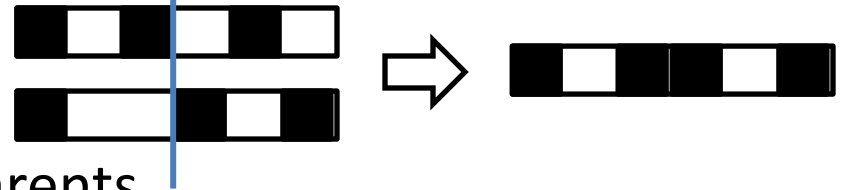
# GA - Parent Selection

- Many methods: roulette wheel, tournament, elitism, etc.
- Elitism selection
  - Keep the best few chromosomes in the population
  - Can perform along with roulette wheel or tournament selection to prevent the solution from getting worse

# GA - Crossover

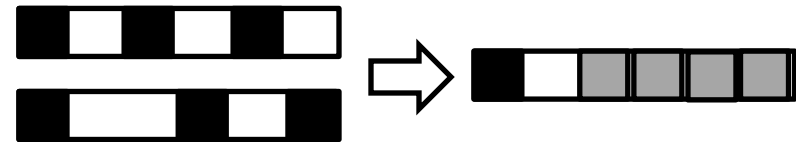
Crossover is used to propagate favorable genes through generations

- **Pure** (for binary chromosome):



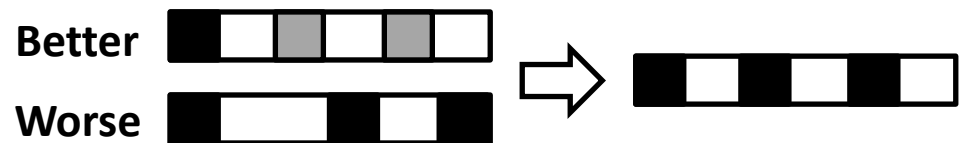
Piecewise combination of two parents

- **Arithmetic** (for real chromosome):



Creates linear interpolation of two parents

- **Heuristic**: Creates linear extrapolation of two parents in direction of better parent



**The choice of crossover scheme is case dependent.**

# GA - Mutation

Many mutations are potentially good...

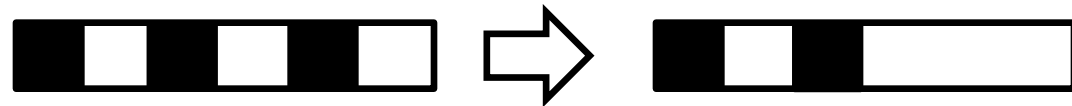


...but too much mutation can reduce fit designs



# GA Mutation Methods

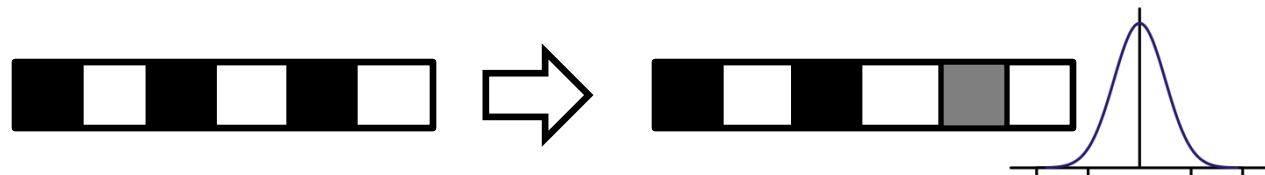
**Boundary:** Set one variable equal to its upper or lower bound



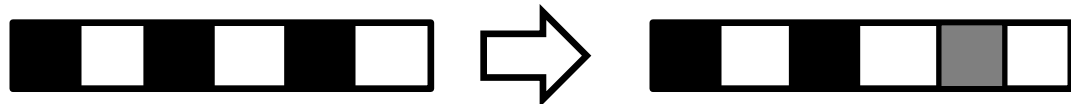
**Uniform:** Set one variable equal to a uniform random number



**Non-uniform:** Set one variable equal to a non-uniform random number



**Incremental:** Increments one variable a random amount



# GA - Termination Criteria

Fixed number of generations

Run out of time

Highest ranking solution reached plateau over last  $K$  iterations



# GA – Pros/Cons

- **Advantages**

- Draws from a large body of designs: global search
- Good performance on combinatorial problems

- **Disadvantages**

- Difficulty balancing size of population/number of generations and overall time
- Genetic operators may not create better designs
- Not necessarily good at fine-tuning a design

# Covariance Matrix Adaptation Evolution Strategy

CMA-ES **learns a covariance matrix** during the evolution, similar to approximating the inverse Hessian in classical methods.

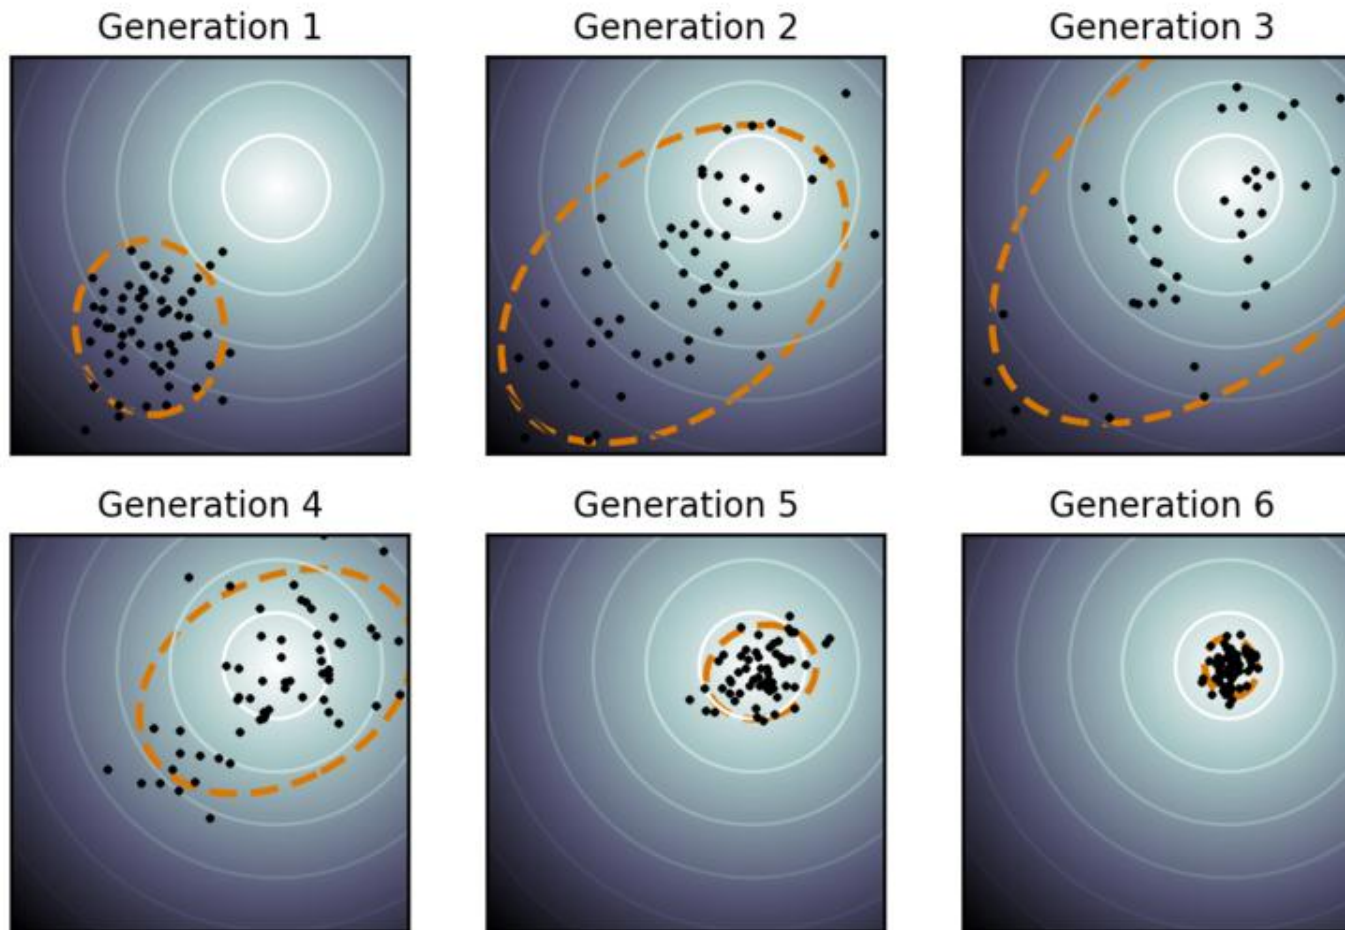


Figure from:  
[http://en.wikipedia.org/wiki/File:Concept\\_of\\_directional\\_optimization\\_in\\_CMA-ES\\_algorithm.png](http://en.wikipedia.org/wiki/File:Concept_of_directional_optimization_in_CMA-ES_algorithm.png)

# CMA-ES Outline

1. Initialize a distribution
2. While *not terminate*
  1. Order the population according to fitness
  2. Pick the top  $\mu$  samples (“good samples”)
  3. Update the mean of good samples
  4. Update the correlation and variance of the distribution
  5. Draw a new population

The updated mean and covariance maximizes the likelihood of “good samples”.

# CMA-ES Pros/Cons

## Pros

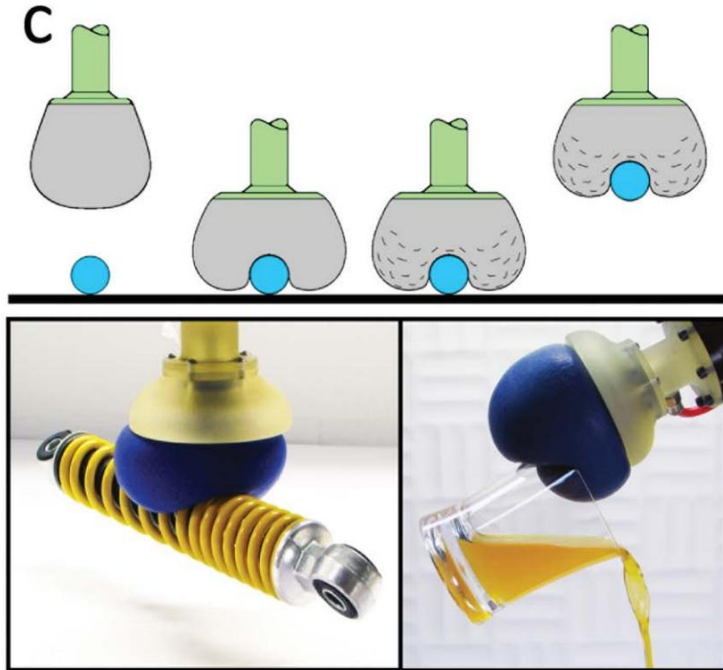
1. Learns a problem specific structure (e.g., better than GA, SA or DIRECT)
2. Suitable for non-convex, non-separable, ill-conditioned, multi-modal or noisy objective functions
3. Suitable for parallel computing
4. Suitable for problems that cannot be solved with a small number of function evaluations ( $<10 \times$  problem size)

## Cons

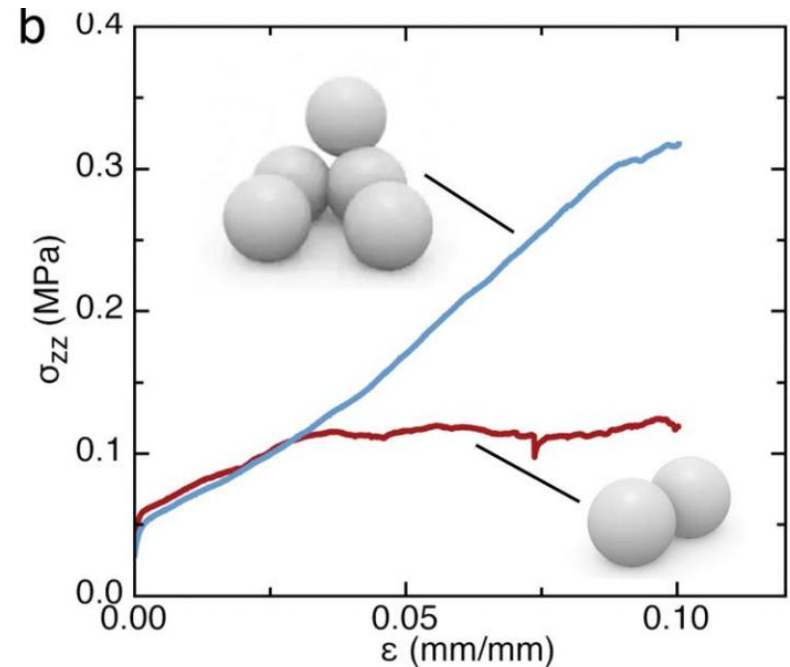
1. Learning does not involve fitness (only ranking)
2. Worse than response surface methods when problem size is small
3. Does not perform well on separable functions

# CMA-ES Applications

## Granular Material design



## Jamming of granular material



## Design material property using evolution strategy

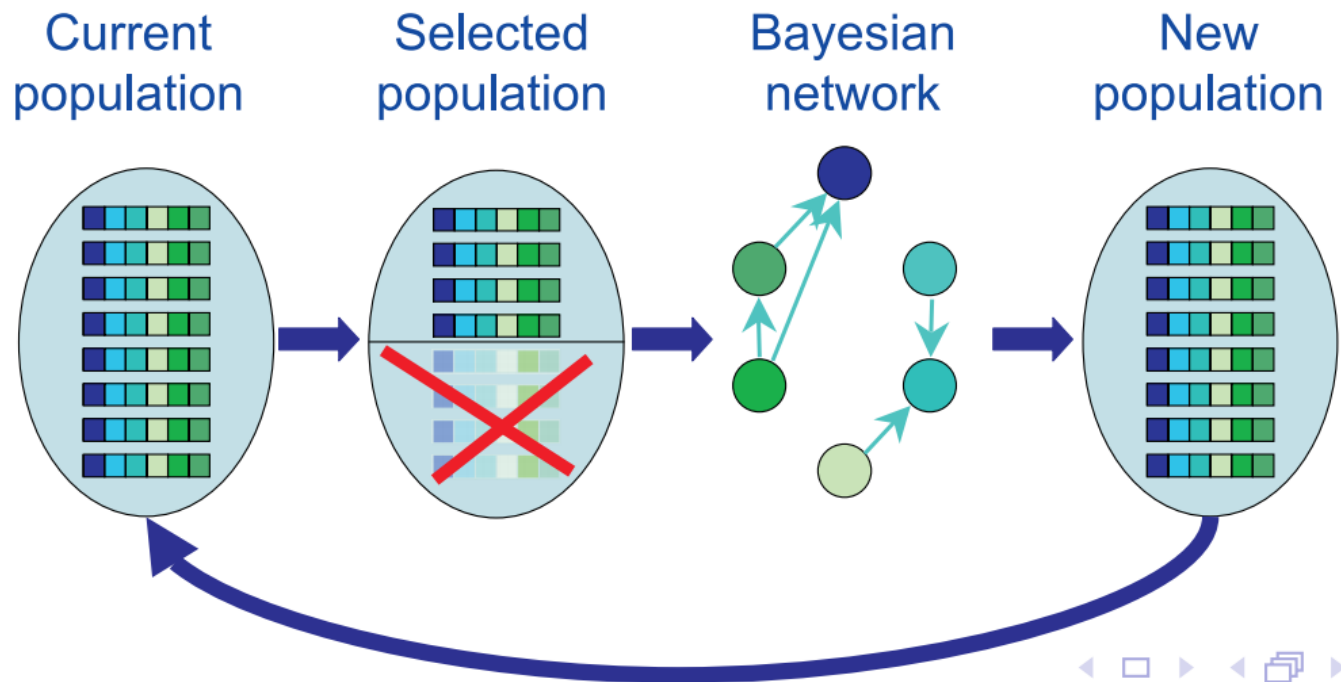
Figures from:  
[http://jfi.uchicago.edu/~jaeger/group/JaegerGroupPapers/granular/Toward\\_jamming\\_by\\_design.pdf](http://jfi.uchicago.edu/~jaeger/group/JaegerGroupPapers/granular/Toward_jamming_by_design.pdf)

# The Bayesian Optimization Algorithm

- The idea of Genetic Algorithm is to mix promising “building blocks” to achieve good solutions.
- Traditional GA operations are shown to be inefficient in preserving partial solutions.
- More sophisticated operations were introduced to address this problem.

# The Bayesian Optimization Algorithm

- BOA learns promising solutions (parents) using a Bayesian network and produces children that have similar properties as parents.



M. Hauschild, M. Pelikan, K. Sastry, D.E. Goldberg, *Using Previous Models to Bias Structural Learning in the Hierarchical BOA*

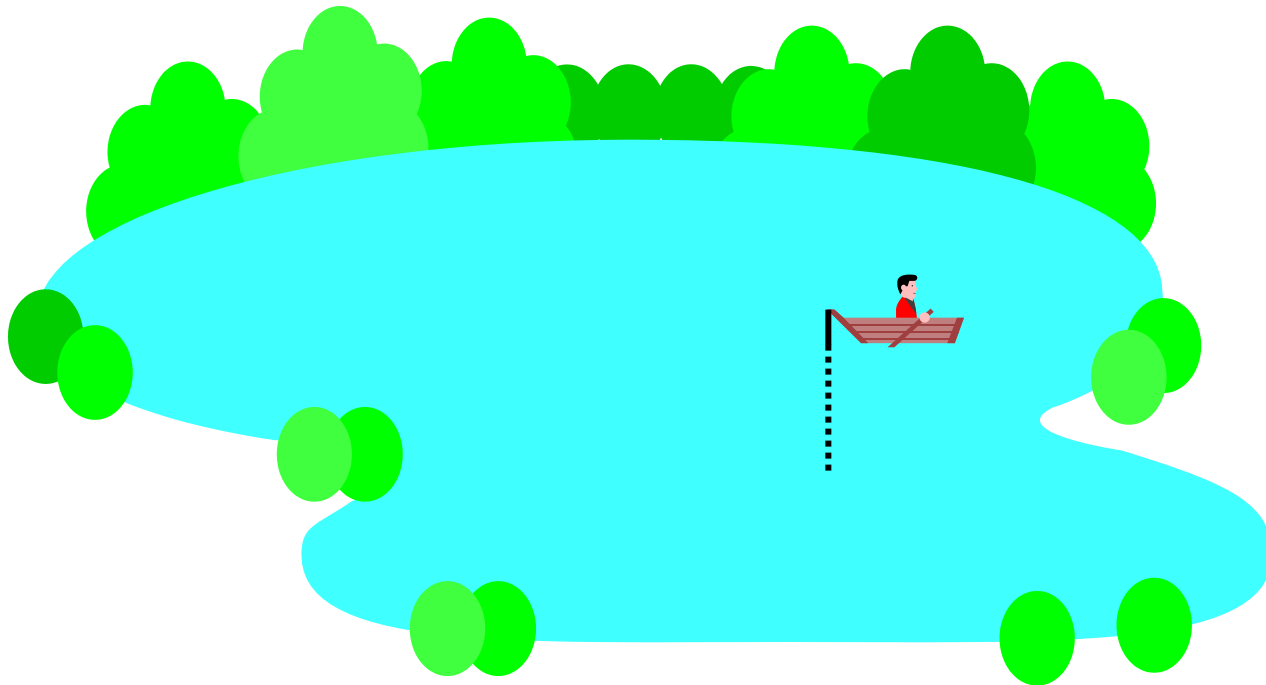
# The Bayesian Optimization Algorithm

- **Advantages:**
  - The learned network preserves good “building blocks”
  - Can handle large decomposable problems more efficiently
- **Disadvantages:**
  - Training networks can be expensive



# EGO – Response Surface

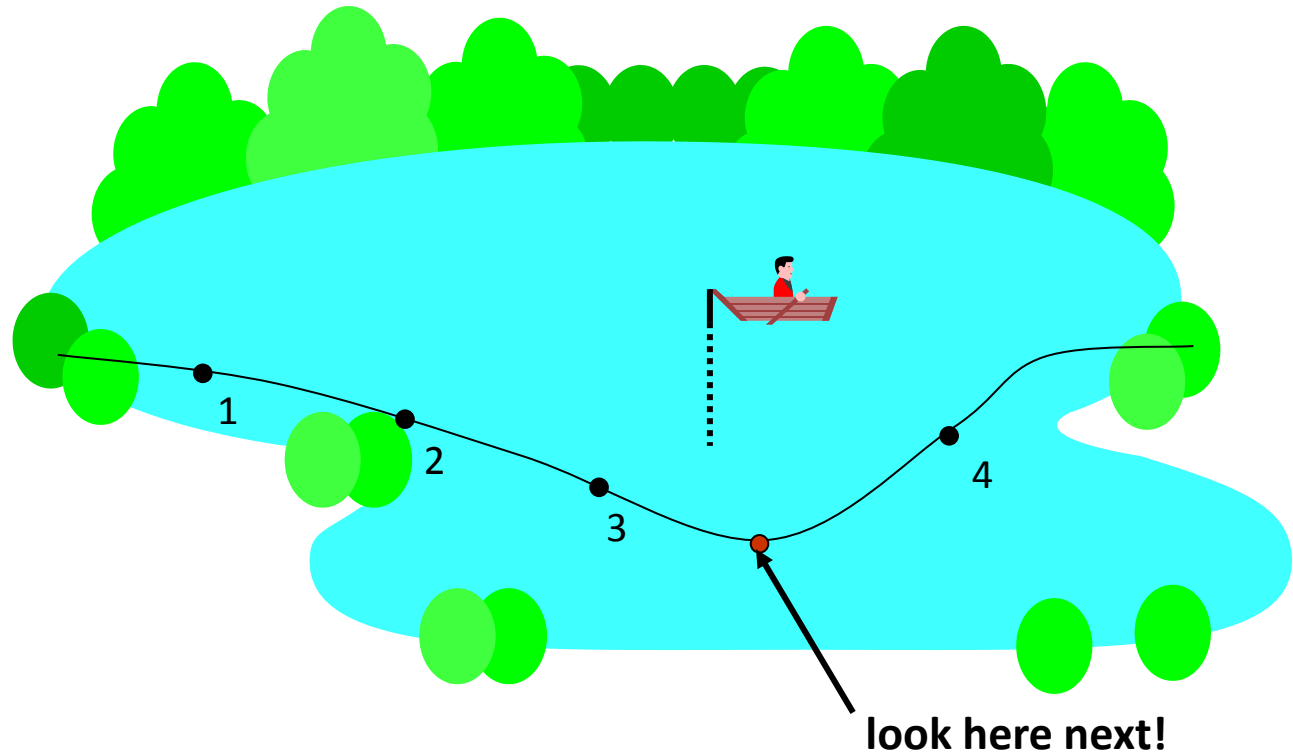
How do you find the deepest part of the lake when you can't see the bottom?



Take a series of depth measurements in strategic locations around the lake.

# EGO – Response Surface

From an initial set of measurements, make a model of the bottom



Use the surrogate model to tell the boat driver where to measure the depth next

# EGO - Kriging

Kriging: A geostatistical techniques to interpolate the elevation of the landscape as a function of the geographic location at an unobserved location from observations of its value at nearby locations.

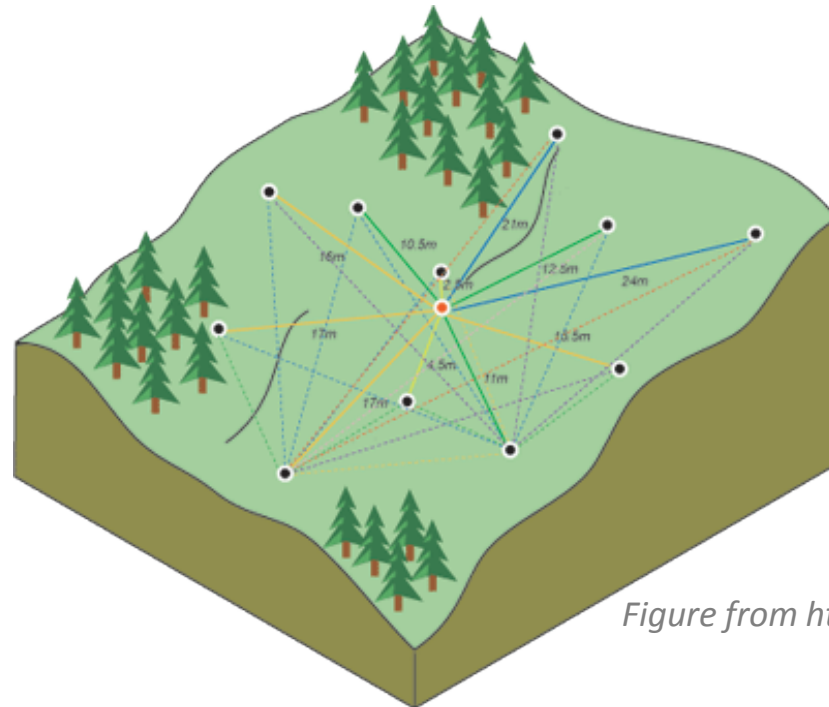


Figure from <http://resources.esri.com>

# EGO – The Merit Function

In each iteration of EGO, we have two functions of  $x$ :

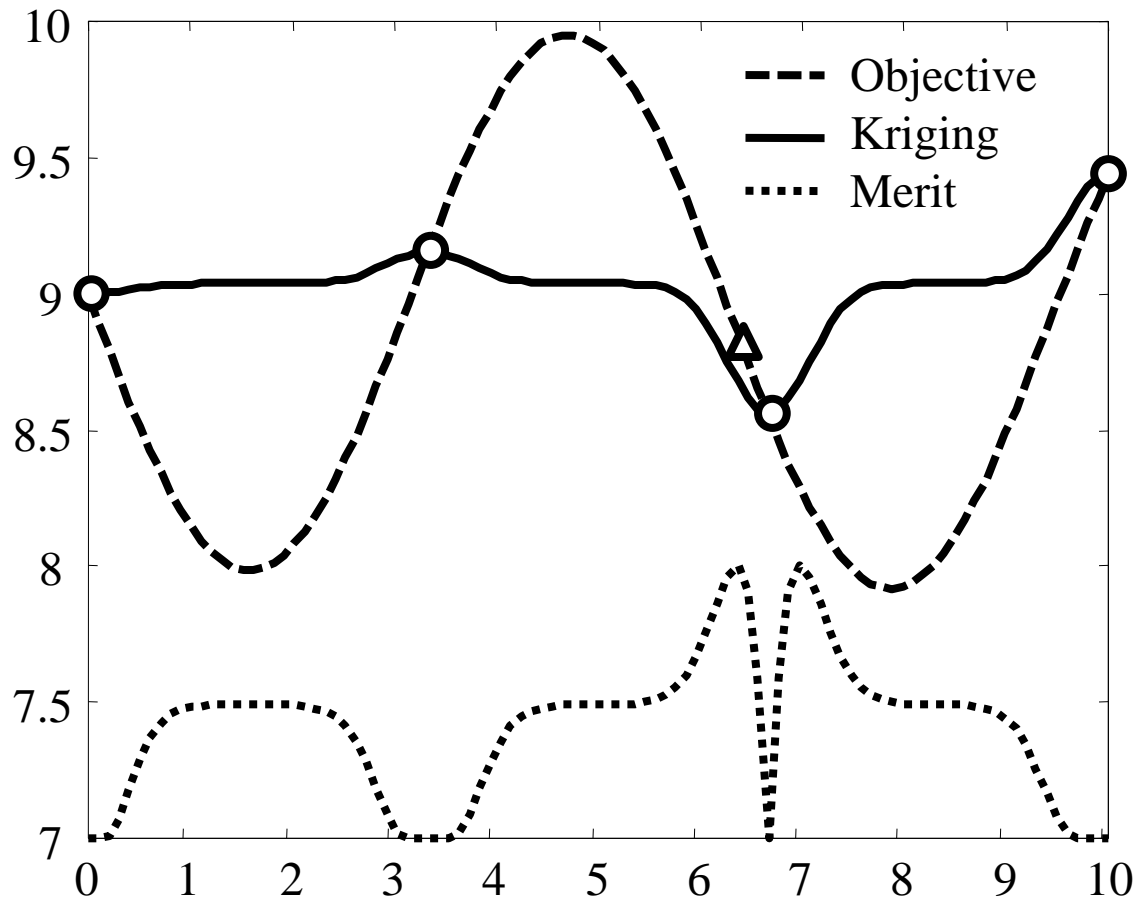
1) the Kriging model  $\hat{y}$ ; 2) the MSE function  $s$ .

The best place to sample next will have low prediction  $\hat{y}$  as well as high uncertainty  $s$ . The merit function reflects the “improvement” of the objective.

$$f_{merit}(x) = (f_{min} - \hat{y})\Phi\left(\frac{f_{min} - \hat{y}}{s}\right) + s\phi\left(\frac{f_{min} - \hat{y}}{s}\right)$$

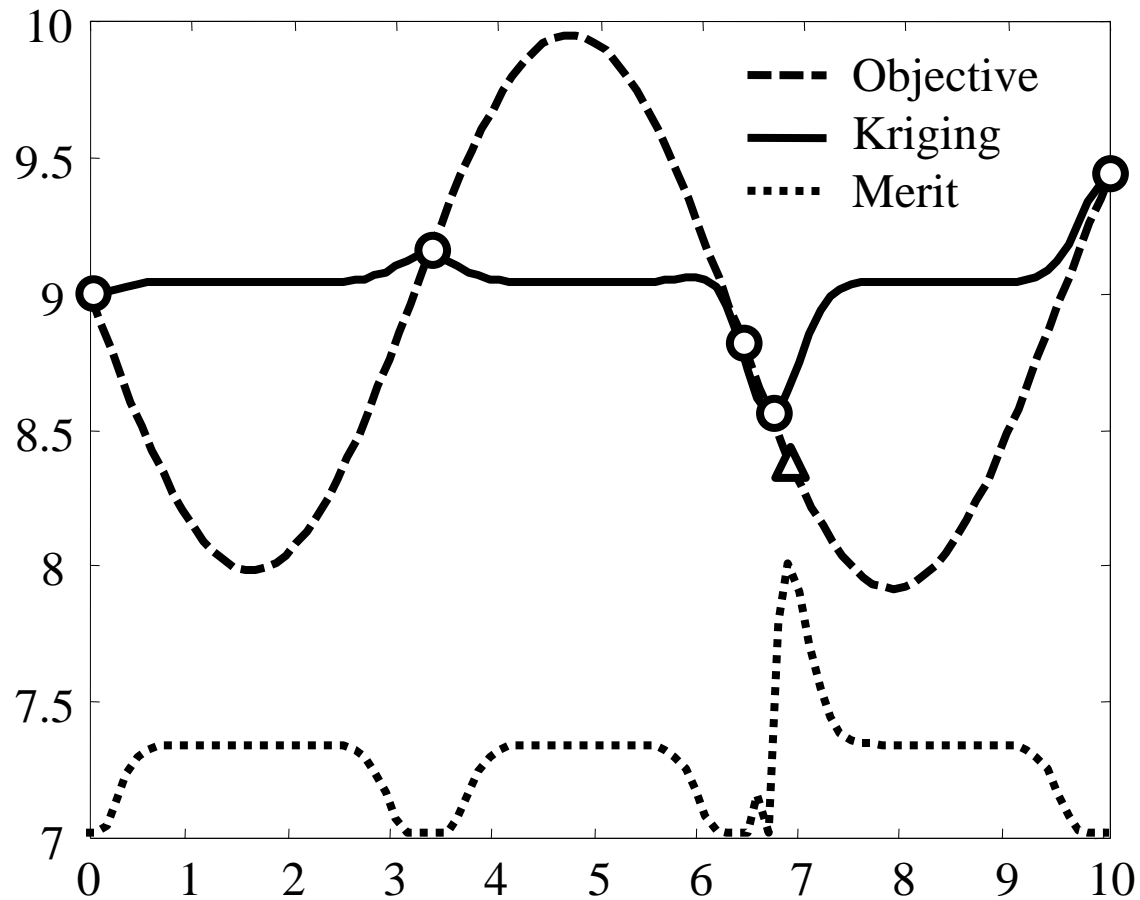
# EGO - Example

Iteration #1



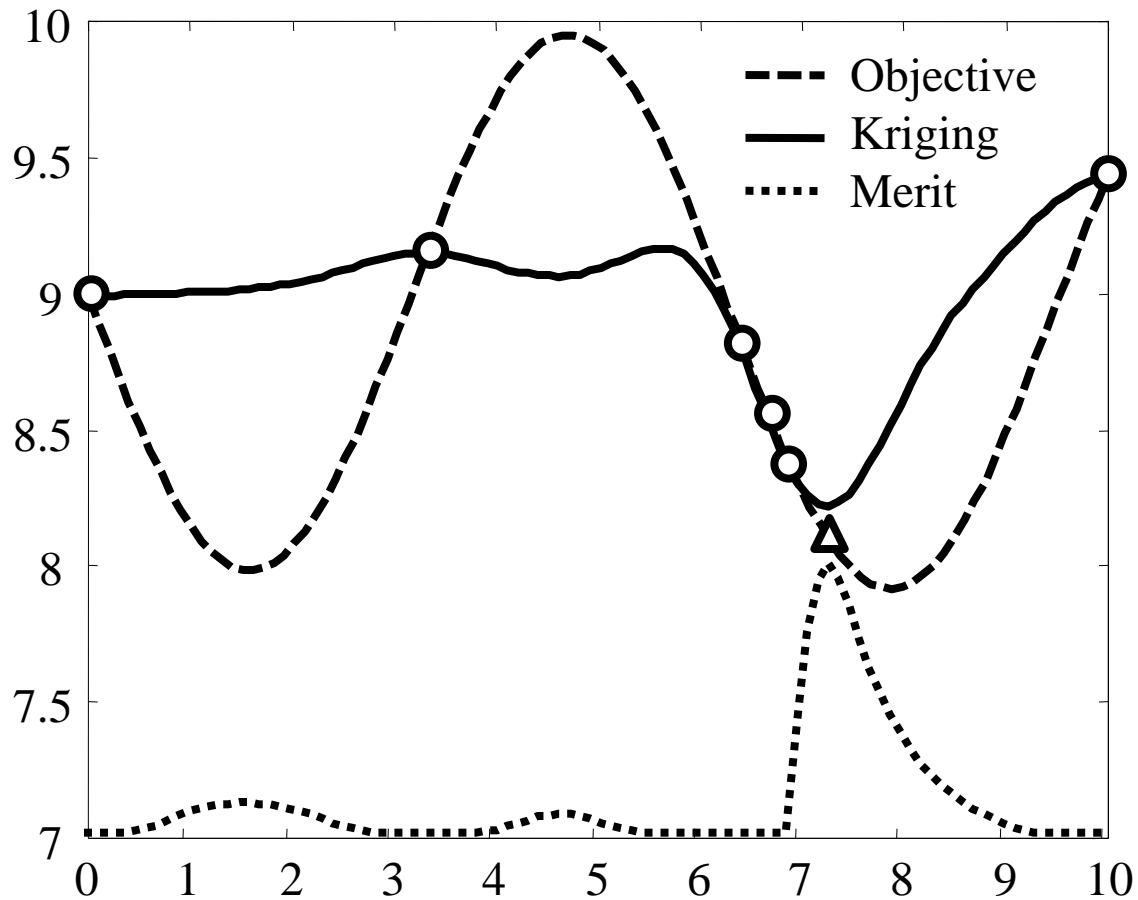
# EGO - Example

Iteration #2



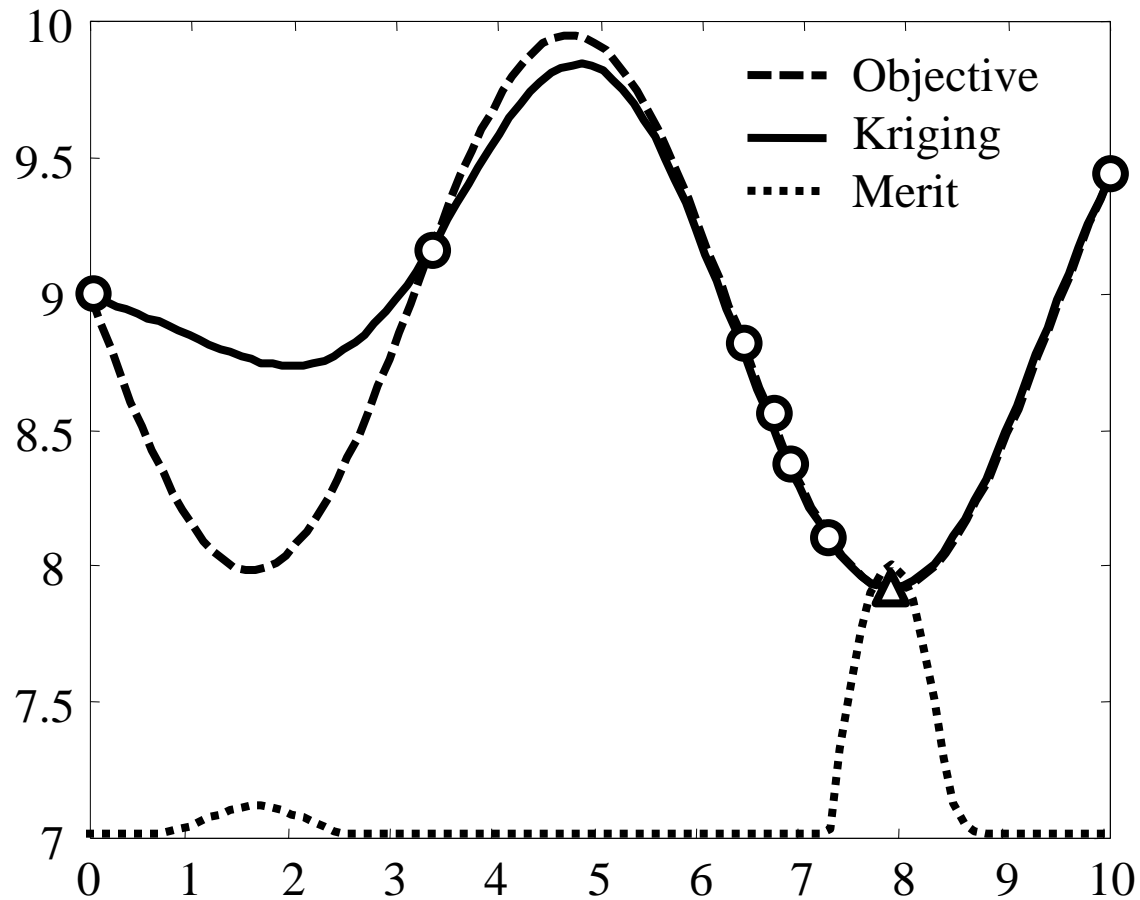
# EGO - Example

Iteration #3



# EGO - Example

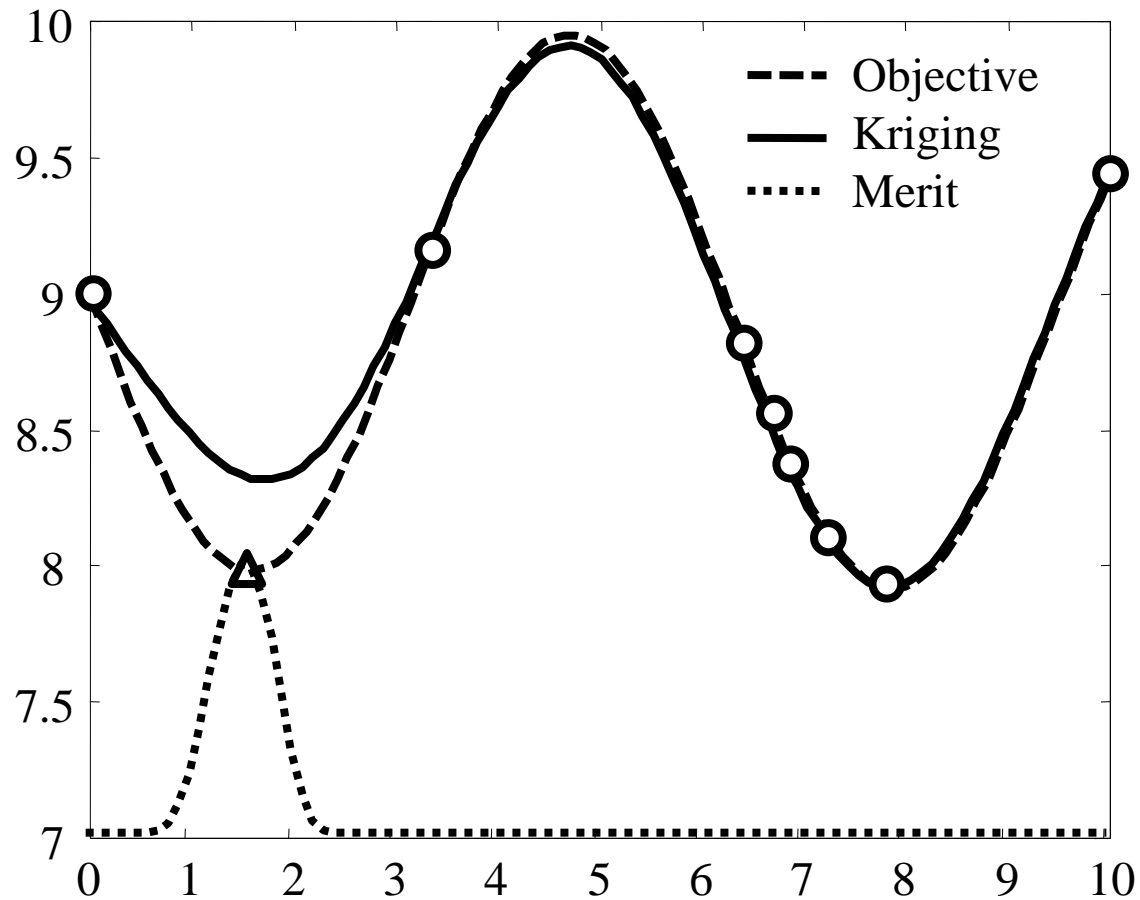
Iteration #4





# EGO - Example

Iteration #5



# EGO – Pros/Cons

- **Advantages**

- Creates surrogate model during search, which is advantageous for expensive functions
- Surrogate model can smooth out noise and discontinuities
- Balances global/local search, similar to DIRECT

- **Disadvantages**

- Difficulty making surrogate model at high dimensions
- Has to create surrogate model for each function, including constraints
- Difficulty optimizing the merit function at high dimensions

# Nelder-Mead Simplex

“Simplex” or “Polytope”

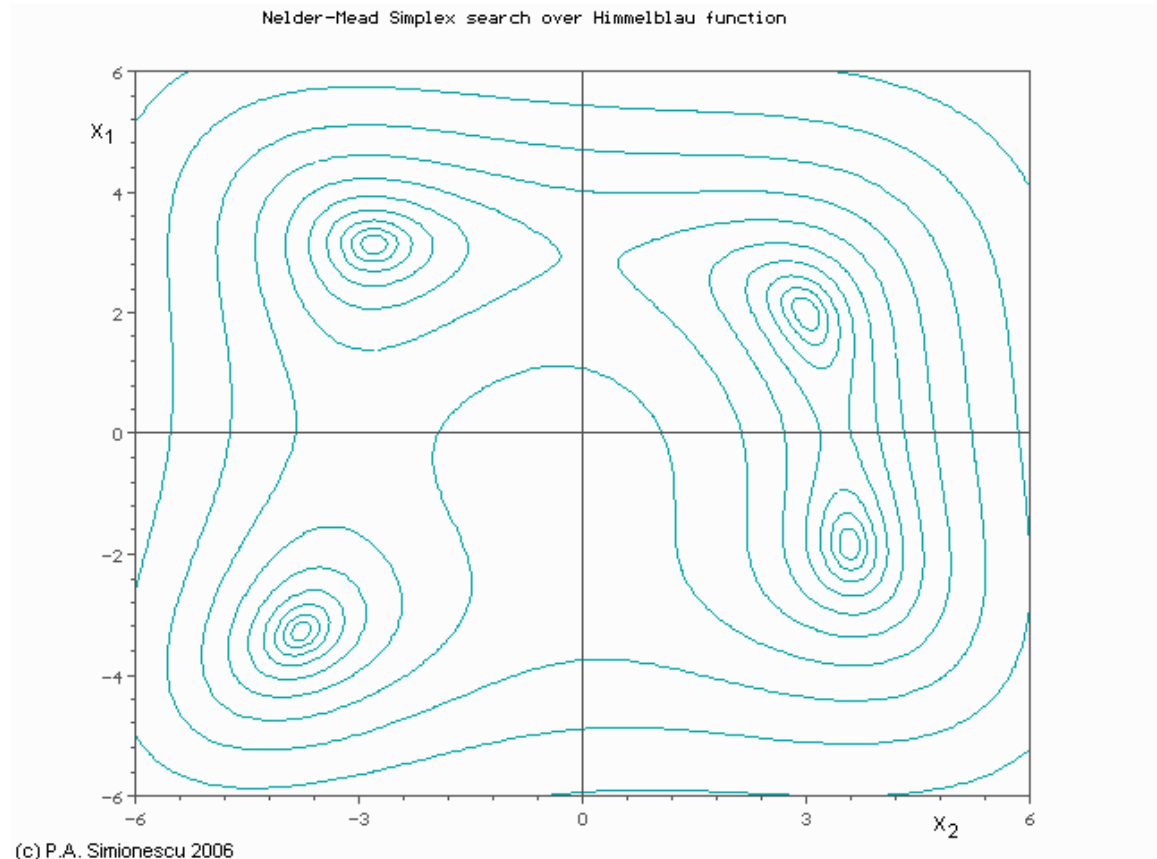
$M+1$  polygon

In this case, 3 vertices

$$\mathbf{x}^{(1)} = (x_1^{(1)}, x_2^{(1)})$$

$$\mathbf{x}^{(2)} = (x_1^{(2)}, x_2^{(2)})$$

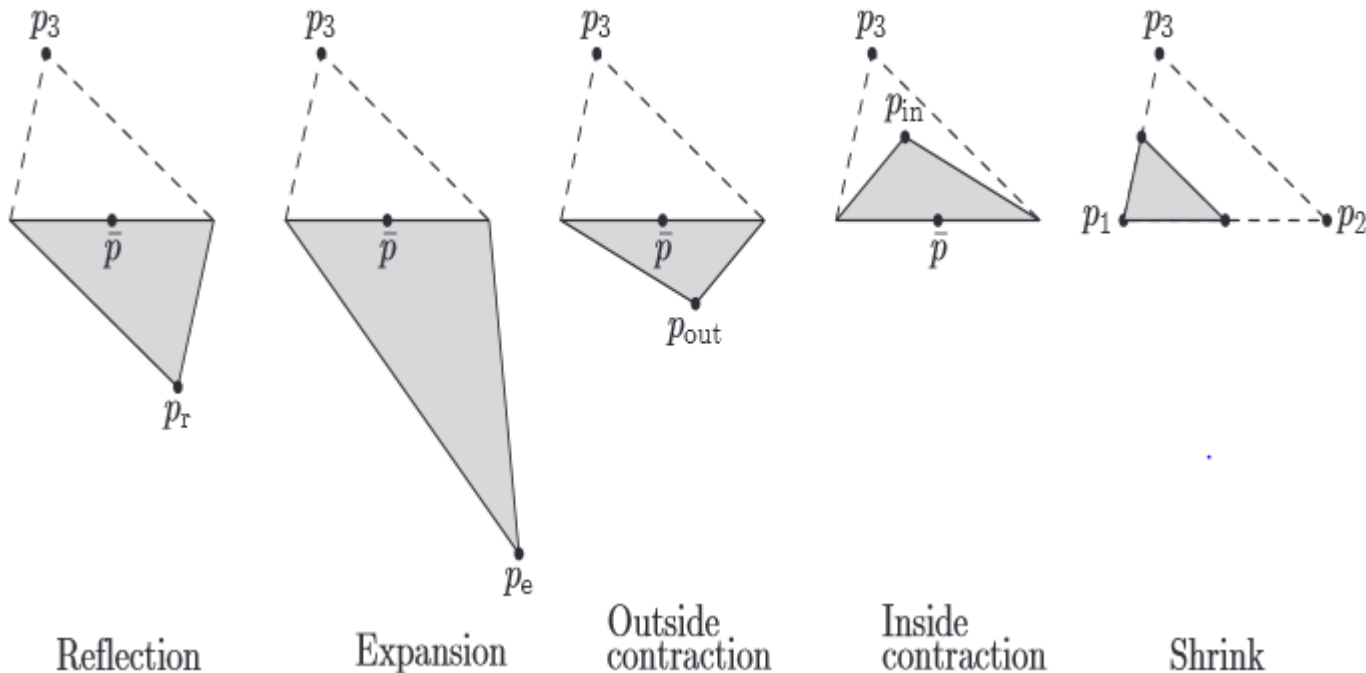
$$\mathbf{x}^{(3)} = (x_1^{(3)}, x_2^{(3)})$$



[https://upload.wikimedia.org/wikipedia/commons/9/96/Nelder\\_Mead2.gif](https://upload.wikimedia.org/wikipedia/commons/9/96/Nelder_Mead2.gif)

The largest  $f(\mathbf{x}^{(i)})$  of the  $i = M + 1$  vertices changed to the centroid of the polytope

# 5 Possible Polytope Changes (in order of when to try them)

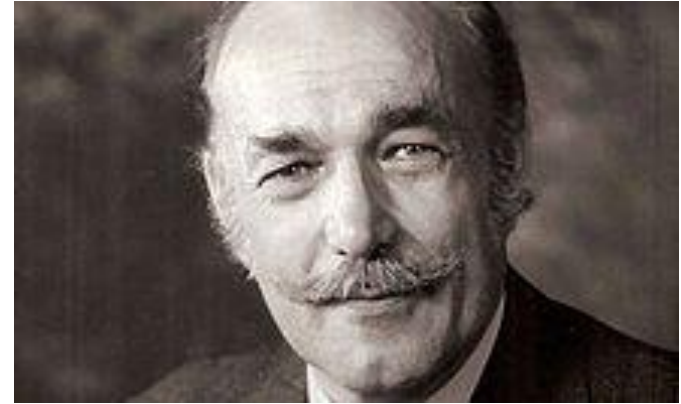


[http://www.math.uiuc.edu/documenta/vol-ismmp/42\\_wright-margaret.pdf](http://www.math.uiuc.edu/documenta/vol-ismmp/42_wright-margaret.pdf)

# Nelder-Mead: Simple, intuitive, and effective

“Mathematicians hate it because you can’t prove convergence; engineers seem to love it because it often works.”

Over 2000 papers cited it in 2012 [1]



**John Nelder**  
National Vegetable  
Research Station

[1] Wright, Margaret H. "Nelder, Mead, and the other simplex method." *Documenta Mathematica* 7 (2010).

# Pros/Cons of Pattern Search

## Advantages

- Easy to implement

- Minimal parameter tuning

## Disadvantages

- Can get stuck easily

- Very dependent on initialization point

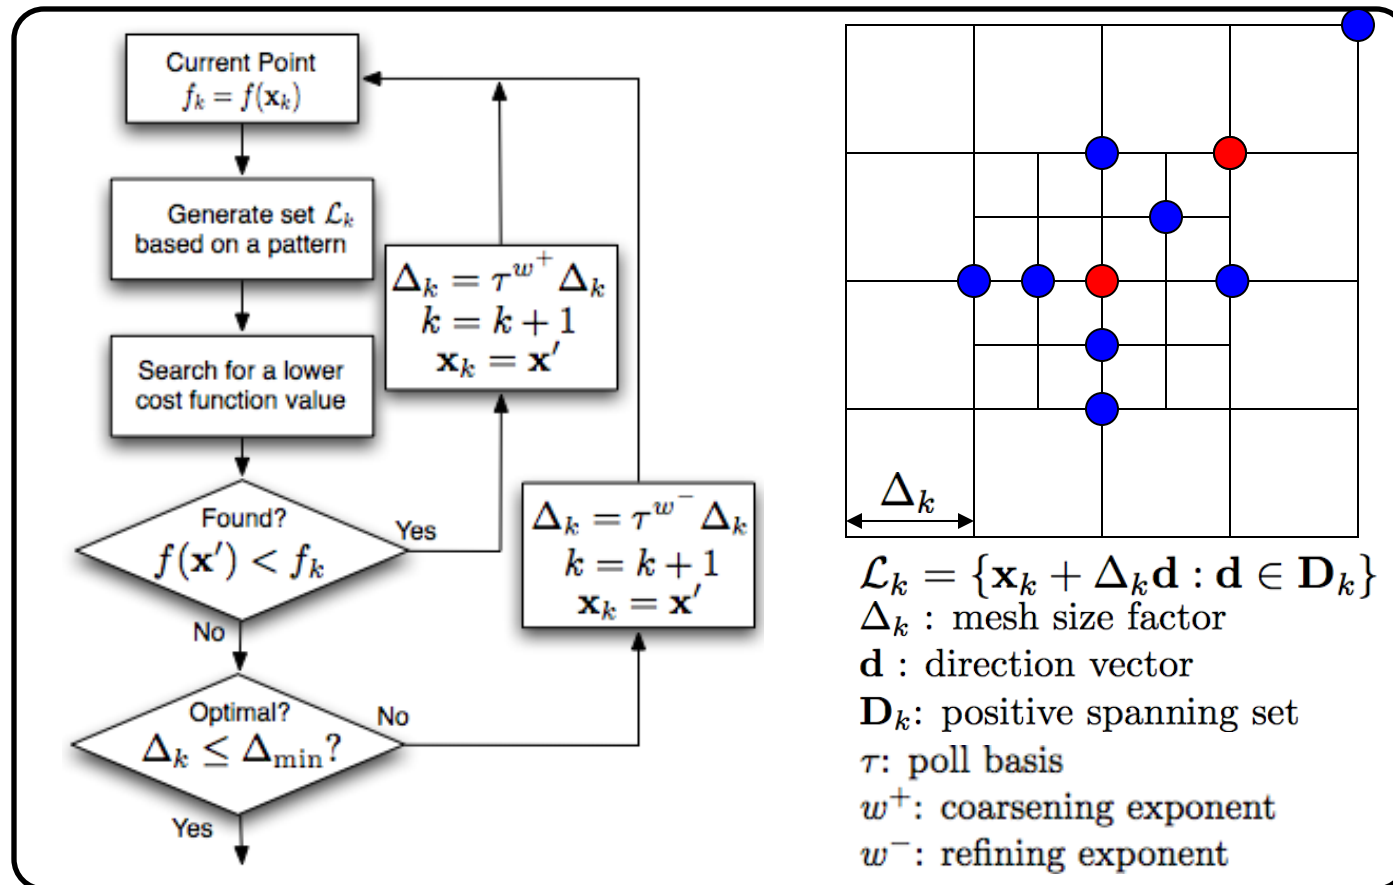
# NOMAD – Overview

- Belongs to Pattern Search
- An implementation of the Mesh-Adaptive Direct Search (MADS) algorithm
- Pattern search method: creates mesh and samples along mesh

# NOMAD – Pattern Search

## Generalized Pattern Search (GPS)

- A number of points around the current point are evaluated
- Best point becomes center point for the next iteration.



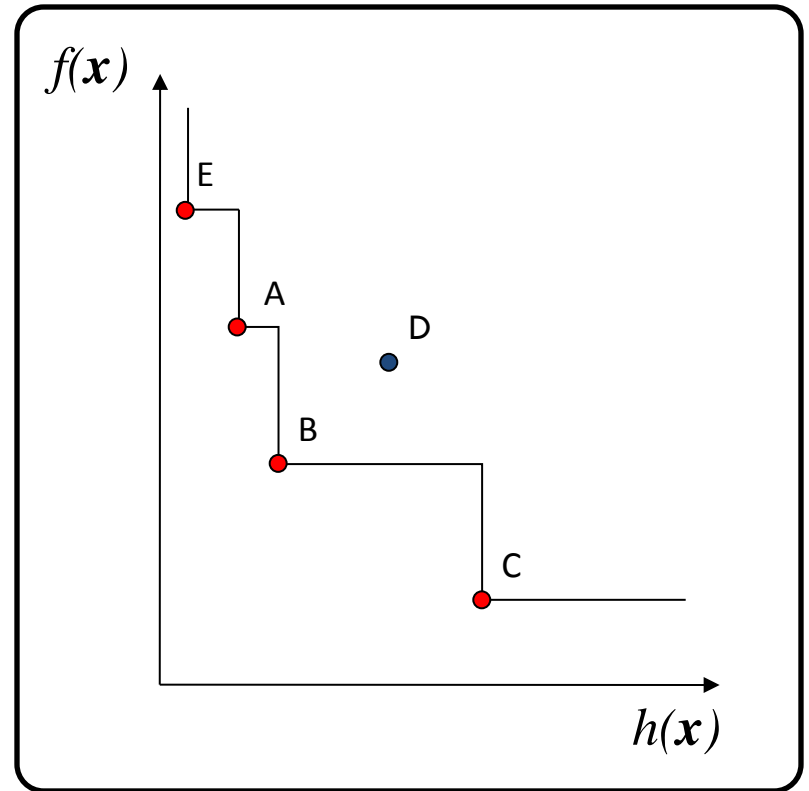


# NOMAD – Constraint

- Bi-objective problem: minimize both the objective function,  $f(x)$ , and an aggregate constraint violation function:

$$h(\bar{x}) = \sum \max \{0, c_i(\bar{x})\}$$

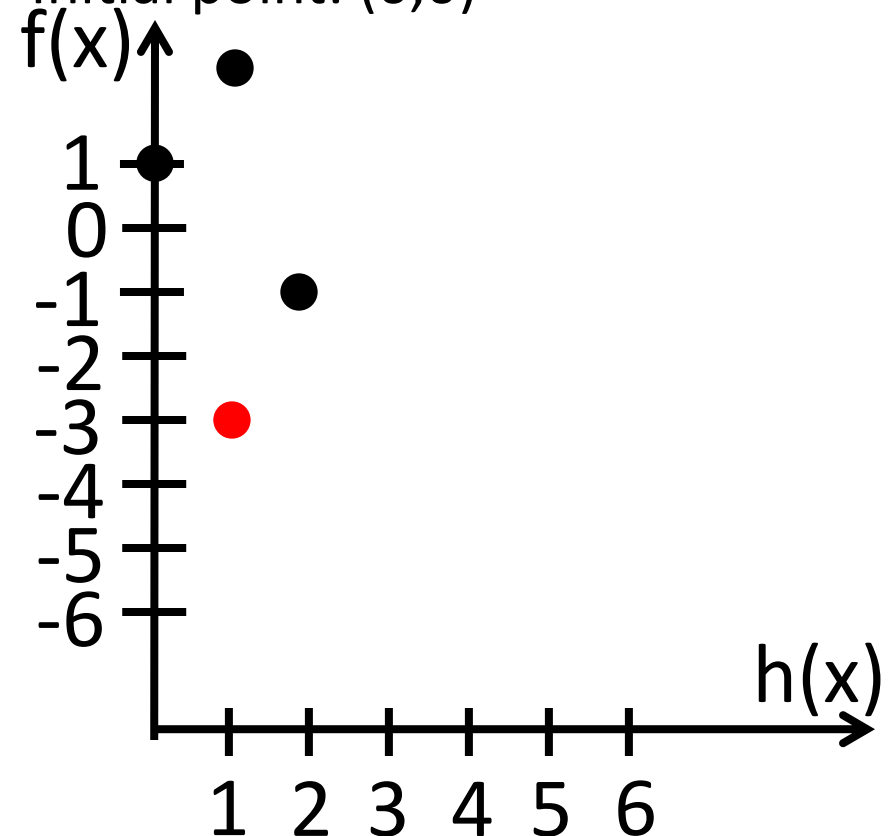
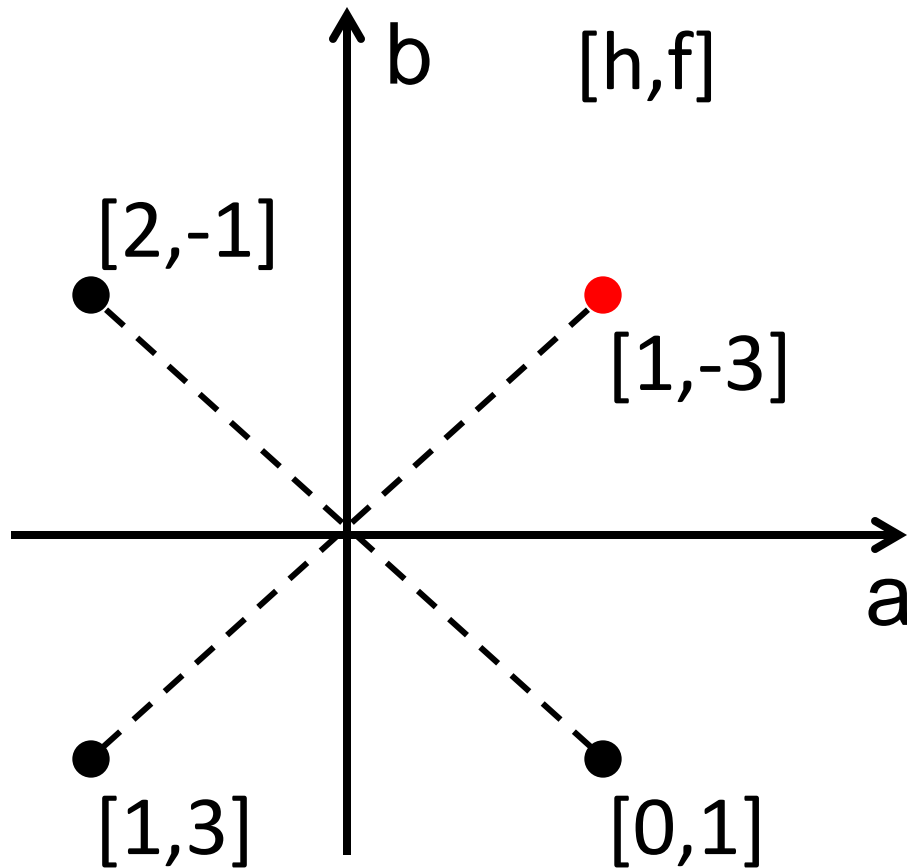
- Chooses Pareto set of Best feasible/Least infeasible points



# GPS– Example

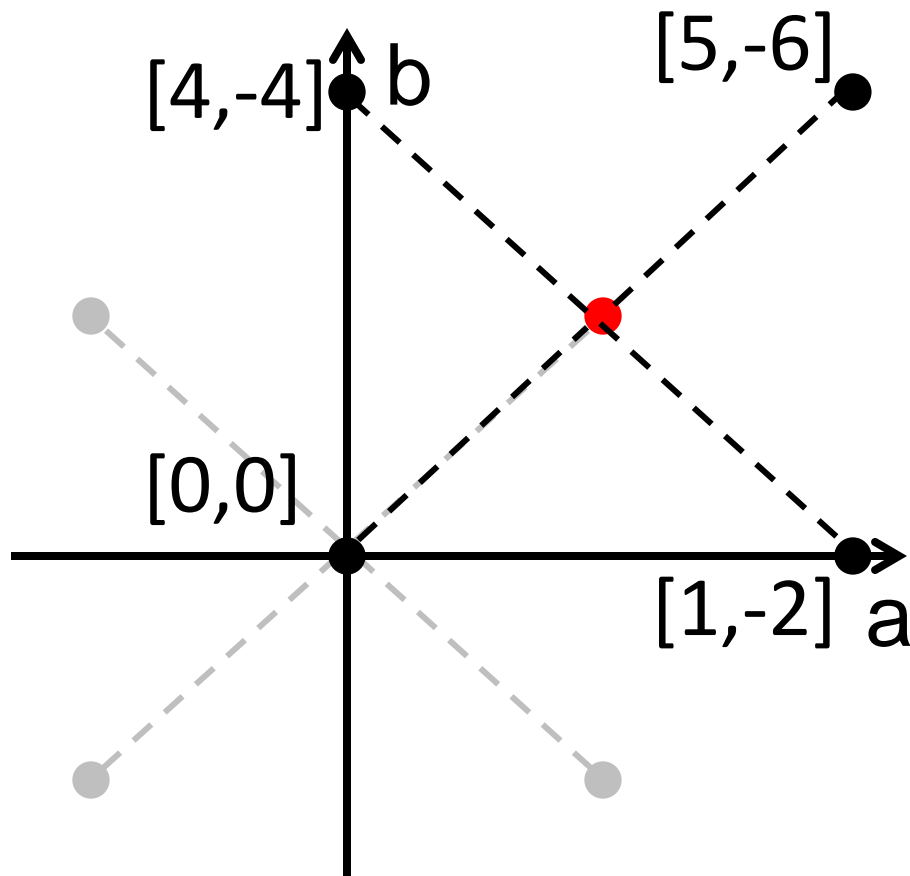
$$\begin{array}{ll} \min_{a,b} & -a - 2b \\ \text{s.t.} & 0 \leq a \leq 1 \\ & b \leq 0 \end{array}$$

GPS, Filter (least infeasible)  
 Directions:  $\pm(1, 1)^T, \pm(1, -1)^T$   
 Initial point:  $(0, 0)^T$

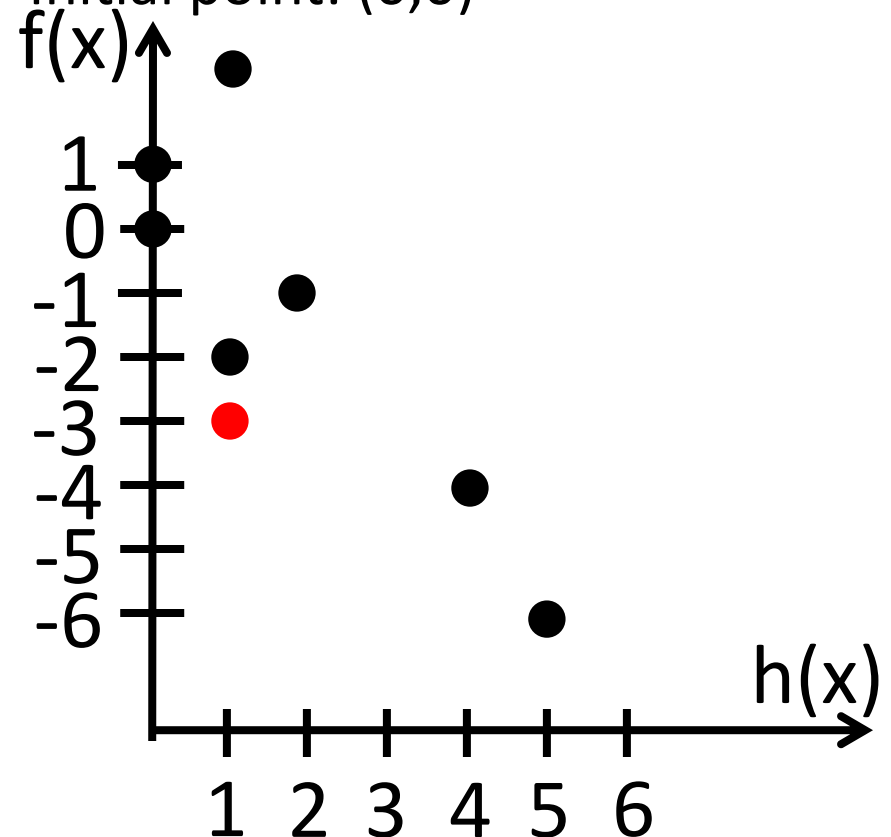


# GPS– Example

$$\begin{array}{ll} \min_{a,b} & -a - 2b \\ \text{s.t.} & 0 \leq a \leq 1 \\ & b \leq 0 \end{array}$$

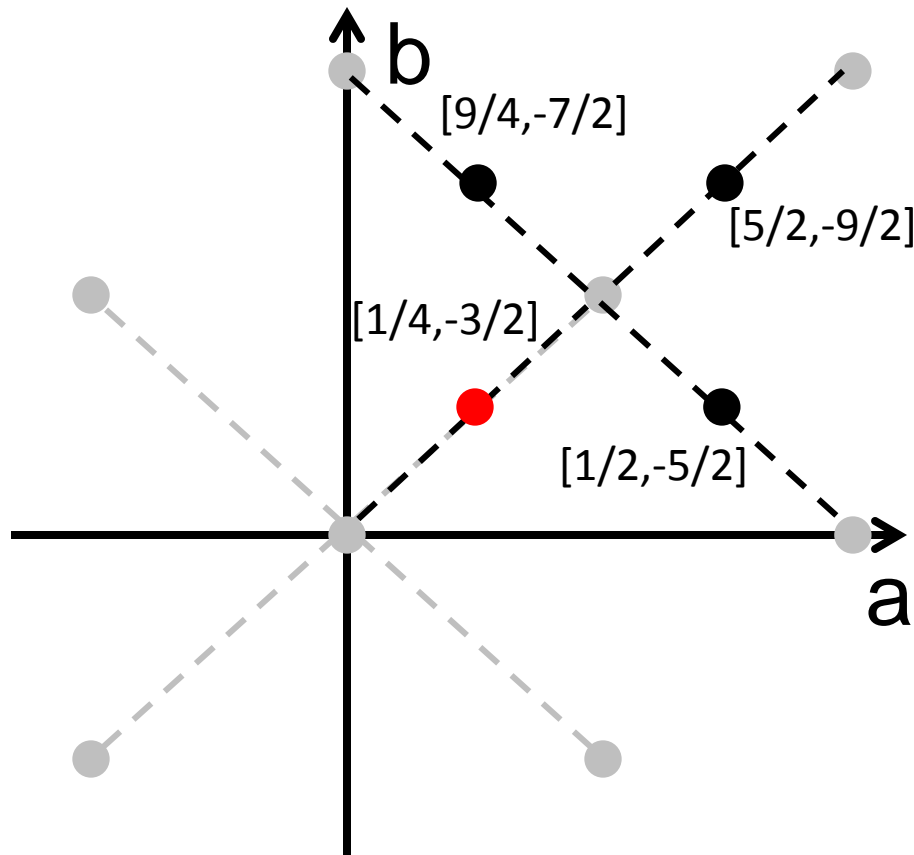


GPS, Filter (least infeasible)  
 Directions:  $\pm(1, 1)^T, \pm(1, -1)^T$   
 Initial point:  $(0,0)^T$

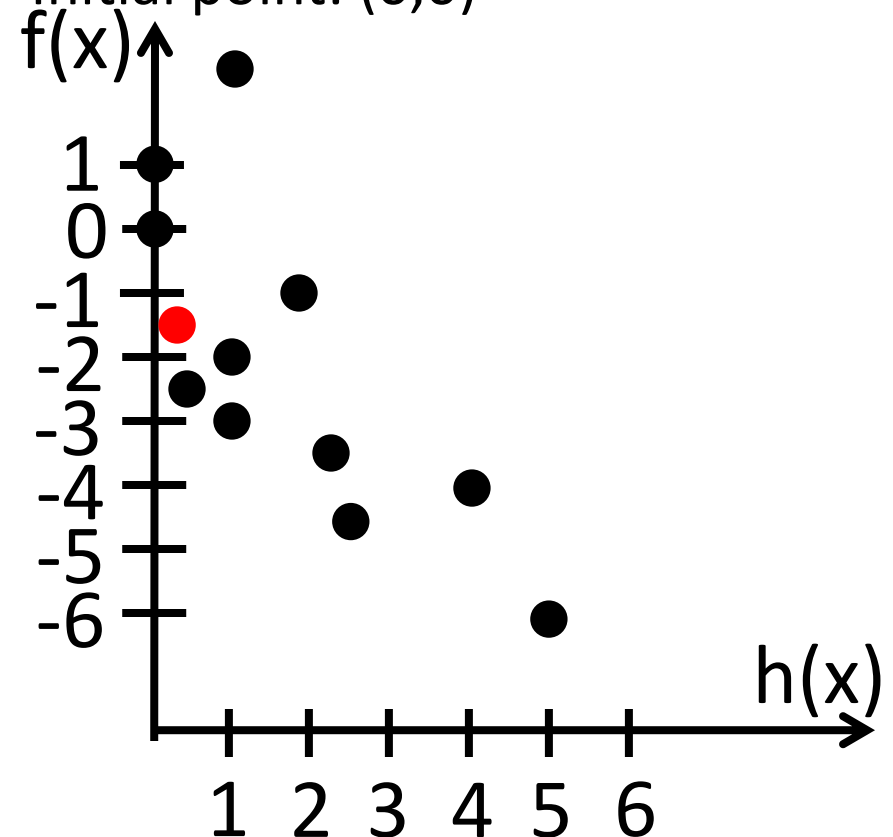


# GPS– Example

$$\begin{array}{ll} \min_{a,b} & -a - 2b \\ \text{s.t.} & 0 \leq a \leq 1 \\ & b \leq 0 \end{array}$$

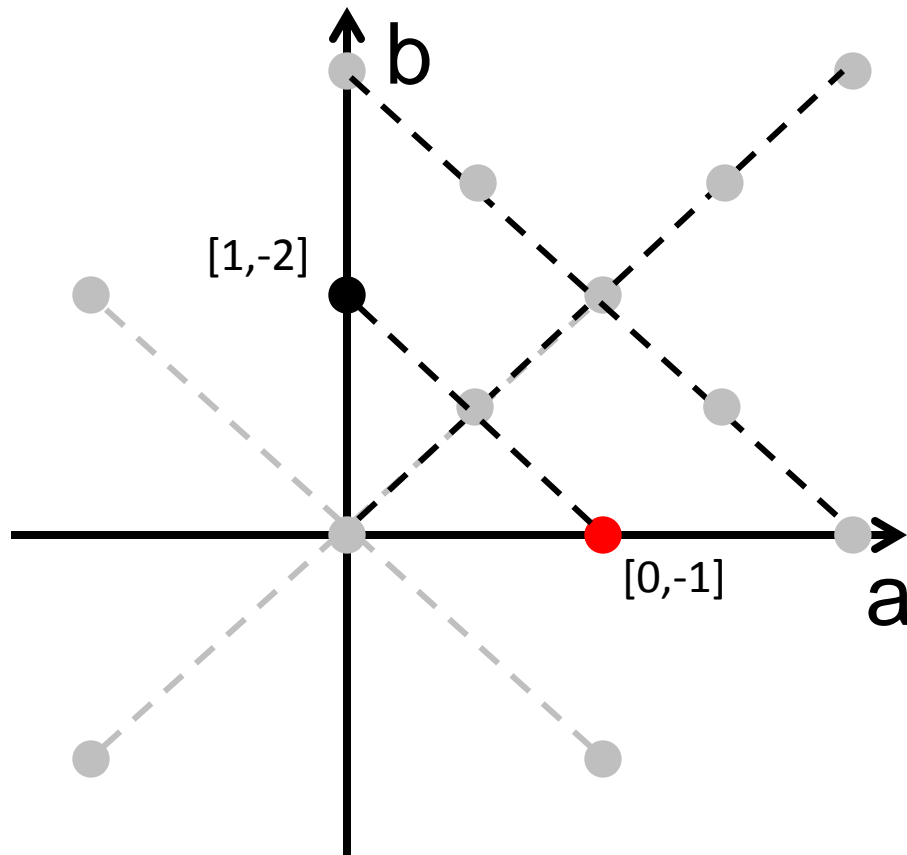


GPS, Filter (least infeasible)  
 Directions:  $\pm(1, 1)^T, \pm(1, -1)^T$   
 Initial point:  $(0, 0)^T$

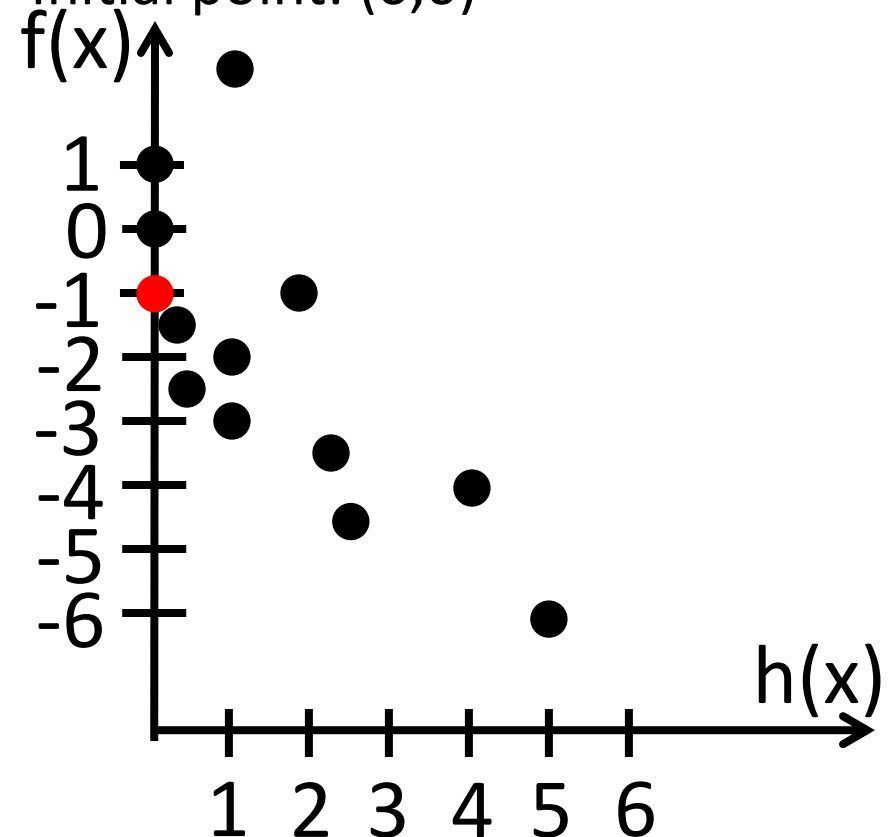


# GPS– Example

$$\begin{array}{ll} \min_{a,b} & -a - 2b \\ \text{s.t.} & 0 \leq a \leq 1 \\ & b \leq 0 \end{array}$$



GPS, Filter (least infeasible)  
 Directions:  $\pm(1, 1)^T, \pm(1, -1)^T$   
 Initial point:  $(0, 0)^T$

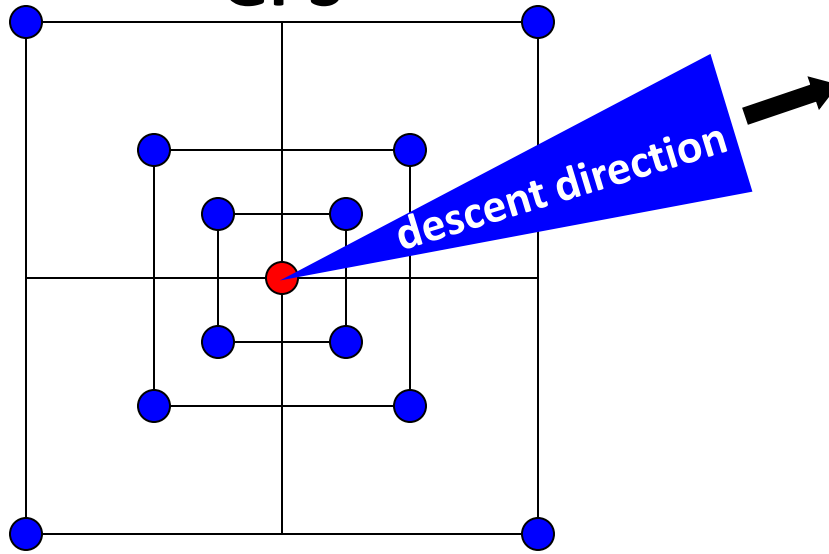


# NOMAD – Pattern Search

- Mesh-Adaptive Direct Search (MADS)
  - GPS shows limitations due to the finite choices of directions
  - MADS removes the GPS restriction by allowing (nearly) infinitely many poll directions
  - Two parameters defining the frame size:  
mesh size  $\Delta_k^m$                       poll size  $\Delta_k^p$
  - mesh size  $\leq$  poll size

# NOMAD – Pattern Search

**GPS**



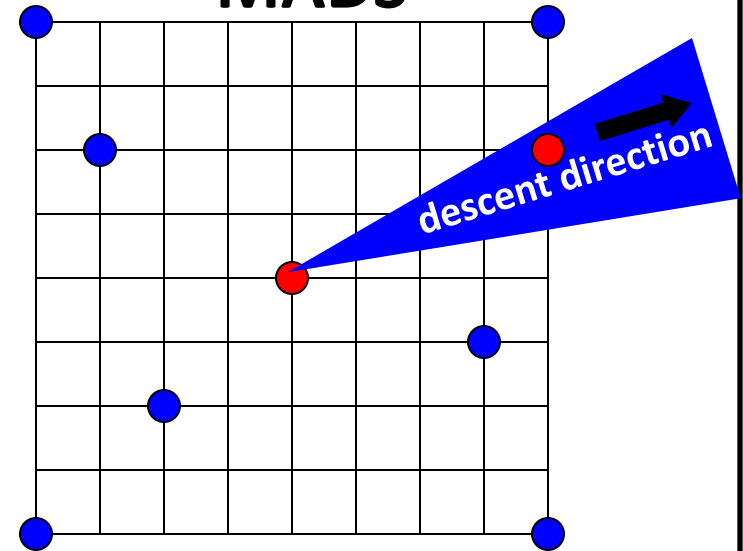
Can't find descent direction  
with finite poll directions

$$\Delta_1^m = \Delta_1^p = 1$$

$$\Delta_2^m = \Delta_2^p = 0.5$$

$$\Delta_3^m = \Delta_3^p = 0.25$$

**MADS**

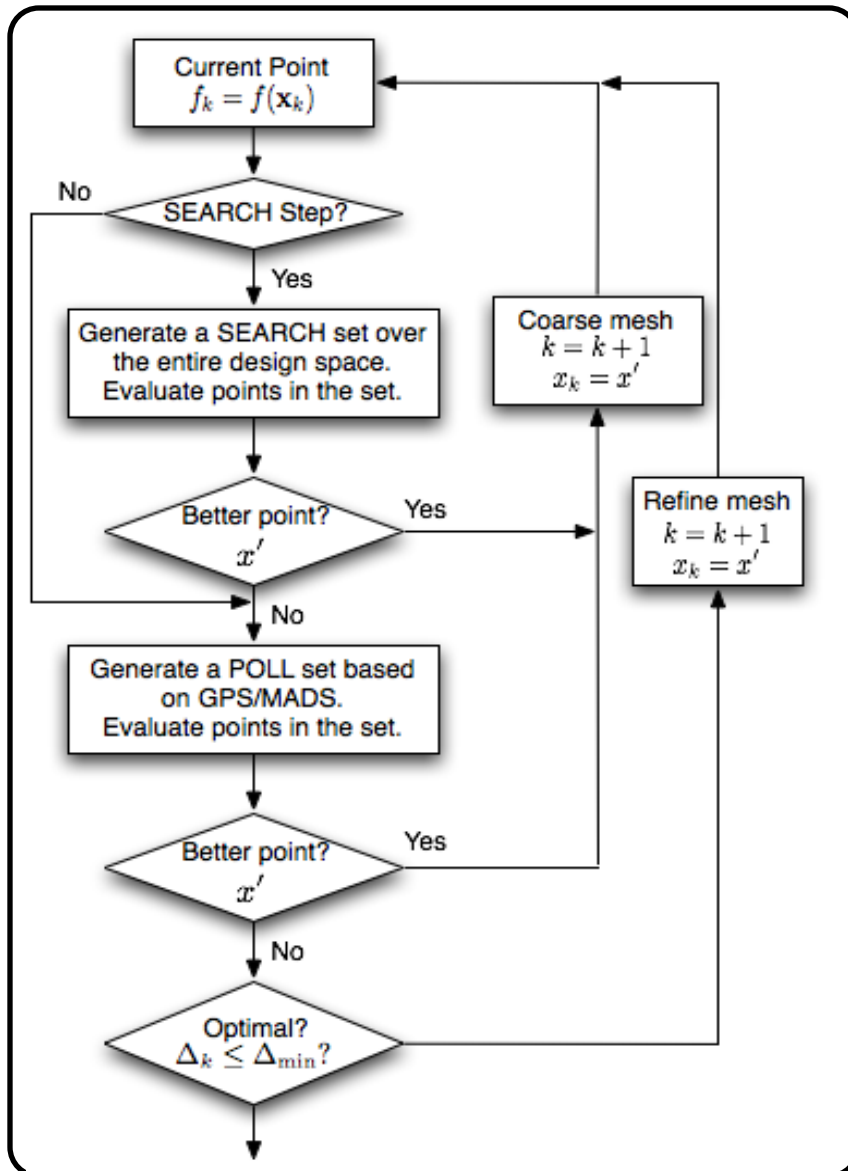


Able to find descent direction  
due to infinitely many poll  
directions

$$\Delta_1^m = 1 \quad \Delta_1^p = 1$$

$$\Delta_2^m = 0.25 \quad \Delta_2^p = 1$$

# NOMAD



- Initial SEARCH step (optional)
  - Random search
  - Genetic algorithm
  - Latin hypercube
  - Orthogonal array
  - Etc.
- POLL step (MADS/GPS)
- Termination criteria based on mesh size



# NOMAD – Pros/Cons

## ■ Advantages

- Can use discrete and categorical variables
- Can integrate other algorithms (e.g. DIRECT) as part of search
- Good combination of Global/Local searching
- Can use gradient information, if available

## ■ Disadvantages

- Poll steps can require a large number of function evaluations in higher dimensions (though  $n+1$  is no larger than finite differencing for a gradient algorithm)
- Can terminate early if gets stuck in one area

# Folk Wisdom – General Optimization

Always start with a gradient method

Try “multistart” with gradient method

Then try a number of *easy* to use nongradient methods...

- ...easy if software exists

- ...easy if not many parameters

| Heuristic Name      | Stochastic/<br>Deterministic   | Constraint Handling | Termination Criteria              | Discrete? | Availability            |
|---------------------|--------------------------------|---------------------|-----------------------------------|-----------|-------------------------|
| Simulated Annealing | Stochastic                     | Weighted Penalty    | min. improvement tolerance        | Y         | Matlab, Optimus, iSight |
| Genetic Algorithm   | Stochastic                     | Weighted Penalty    | #generations/<br>fitness change   | Y         | Matlab, iSight          |
| DIRECT              | Deterministic                  | Weighted Penalty    | #function calls                   | Y         | Matlab, Tomlab          |
| EGO                 | Stochastic or<br>Deterministic | Response Surface    | ask Optimus                       | N         | Tomlab, Optimus         |
| NOMAD               | Stochastic or<br>Deterministic | Pareto Set          | min. mesh size<br>#function calls | Y         | Matlab                  |

# Credits

Some material adapted from:

A. Burnap, AE. Bayrak, N. Kang, University of  
Michigan

and

Prof. Kokkalaras, McGill University

*Thank you*